


Disclaimer

- These slides have been prepared by B. Barla Cambazoglu, Director of Applied Science, NTENT Inc.
- Used in several IR tutorials and schools by Ricardo Baeza-Yates and B. Barla Cambazoglu.
- Used with permission, please do not redistribute in any form.
- Please refer to:



Indexing

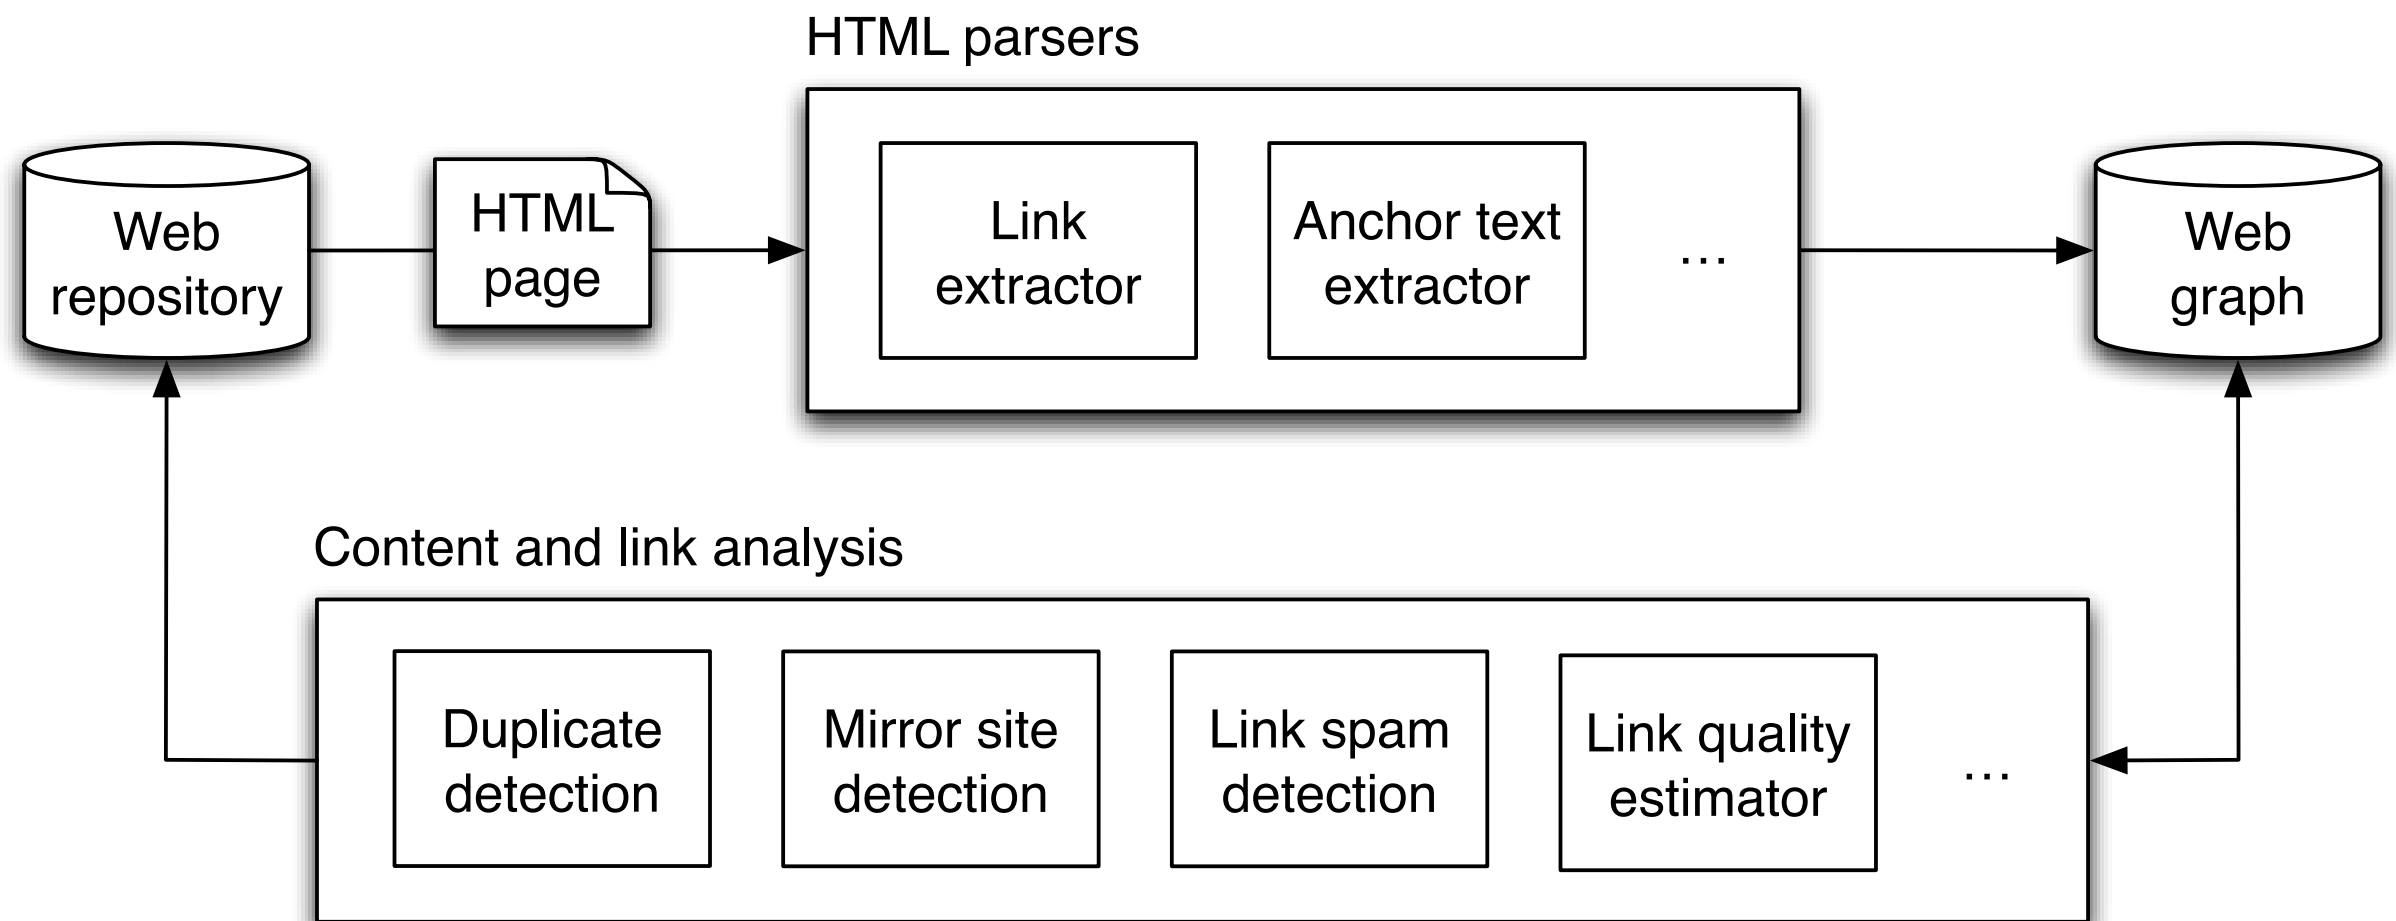
The indexing system performs several tasks

- performs information extraction, filtering, and classification on downloaded web pages
- provides meta-data, metrics, and other kinds of feedback to the crawling and query processing systems
- converts the pages in the web repository into appropriate index structures that facilitate searching the textual content of pages.

Document Processing Pipeline


- A typical indexing system involves various document processing pipelines, each performing different normalization or extraction tasks on web pages.
- Common data structures generated by these pipelines are
 - web graph
 - page attribute file
 - inverted index

Web Graph



Web Graph

- Web graph
 - node: attributes about the page
 - URL
 - inbound/outbound links
 - geographical region
 - language
 - edges: attributes about the links
 - anchor text
- Built at different granularities
 - page-level: duplicate detection
 - host-level: host quality estimation
 - site-level: mirror site detection




Link Analysis

- PageRank: A link analysis algorithm that assigns a weight to each web page indicating its importance.
- Iterative process that converges to a unique solution.
- Weight of a page is proportional to
 - number of inbound links of the page
 - weight of linking pages
- Other algorithms
 - HITS
 - TrustRank

Spam Detection

- Types of spam: link spam, content spam, cloaking/redirection spam, click spam.

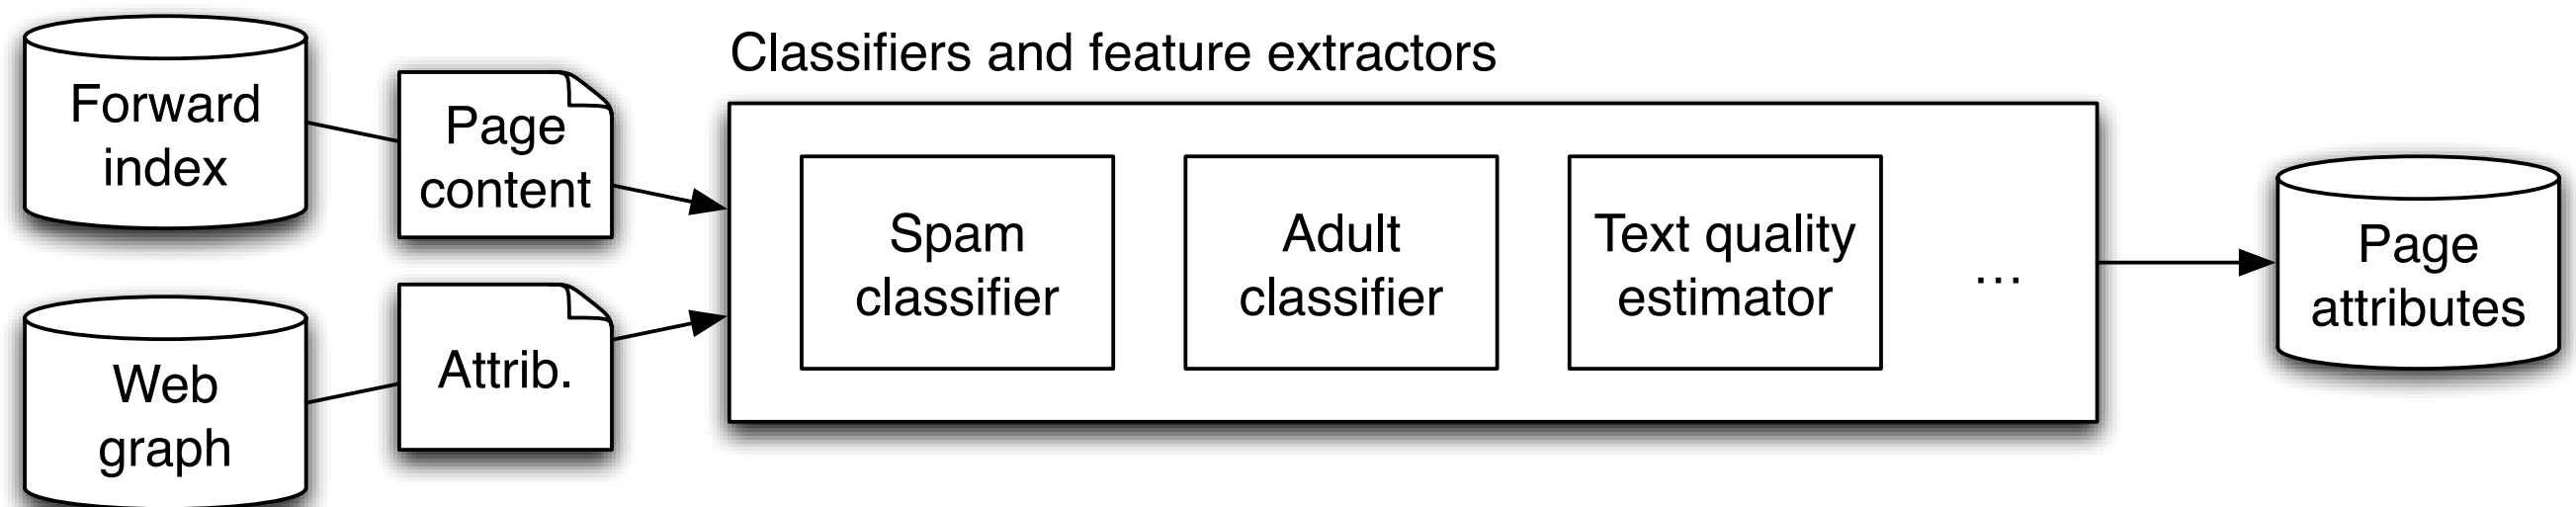


Duplicate Page Detection

- Detecting pages that have duplicate content
 - exact duplicates
 - comparing hash values
 - near duplicates
 - shingles
 - locality sensitive hashing

P1: A B C D E F 79, 189, 44, 14, 99 $H1 = \{14, 44, 79\}$
 └───┘ └───┘ $\longrightarrow J(H1, H2) = 4/6$
P2: A B C X D E F 79, 189, 84, 68, 14, 99 $H2 = \{14, 68, 79\}$

Page Attribute File



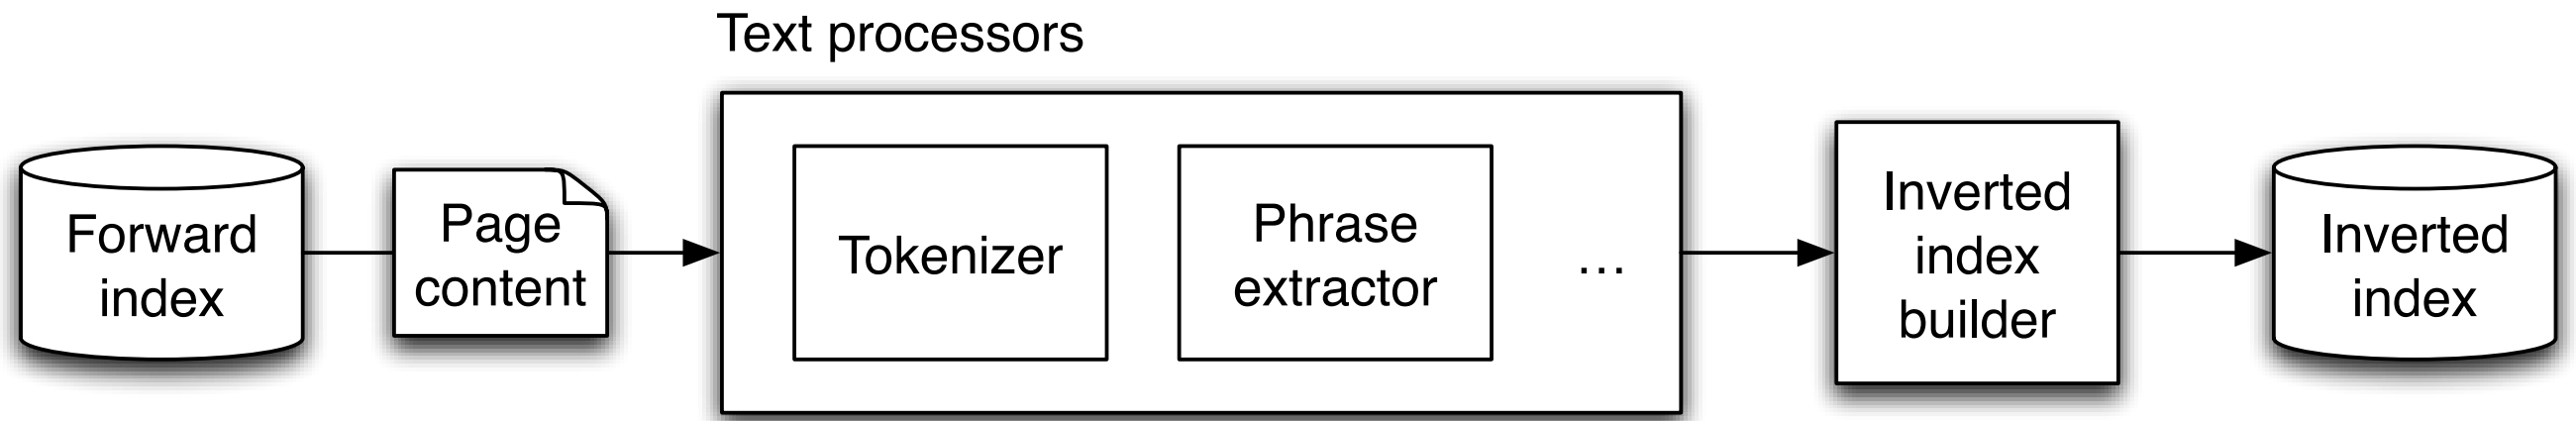
doc id: 4

| | | | | | |
|--|------------|--------|------------|-----|--|
| | 7 | 240 | 1.7 | ... | |
| | term count | length | spam score | | |

Query-Independent Features

| Feature | Source | Description |
|----------------|--------------|--|
| Content spam | Page content | Score indicating the likelihood that the page content is spam |
| Text quality | Page content | Score combining various text quality features (e.g., readability) |
| Link quality | Web graph | Page importance estimated based on page's link structure |
| CTR | Query logs | Observed click-through rate of the page in search results (if available) |
| Dwell time | Query logs | Average time spent by the users on the page |
| Page load time | Web server | Average time it takes to receive the page from the server |
| URL depth | URL string | Number of slashes in the absolute path of the URL |

Inverted Index



- Text processing may involve

- tokenization
- stopword removal
- case conversion
- stemming

- Example

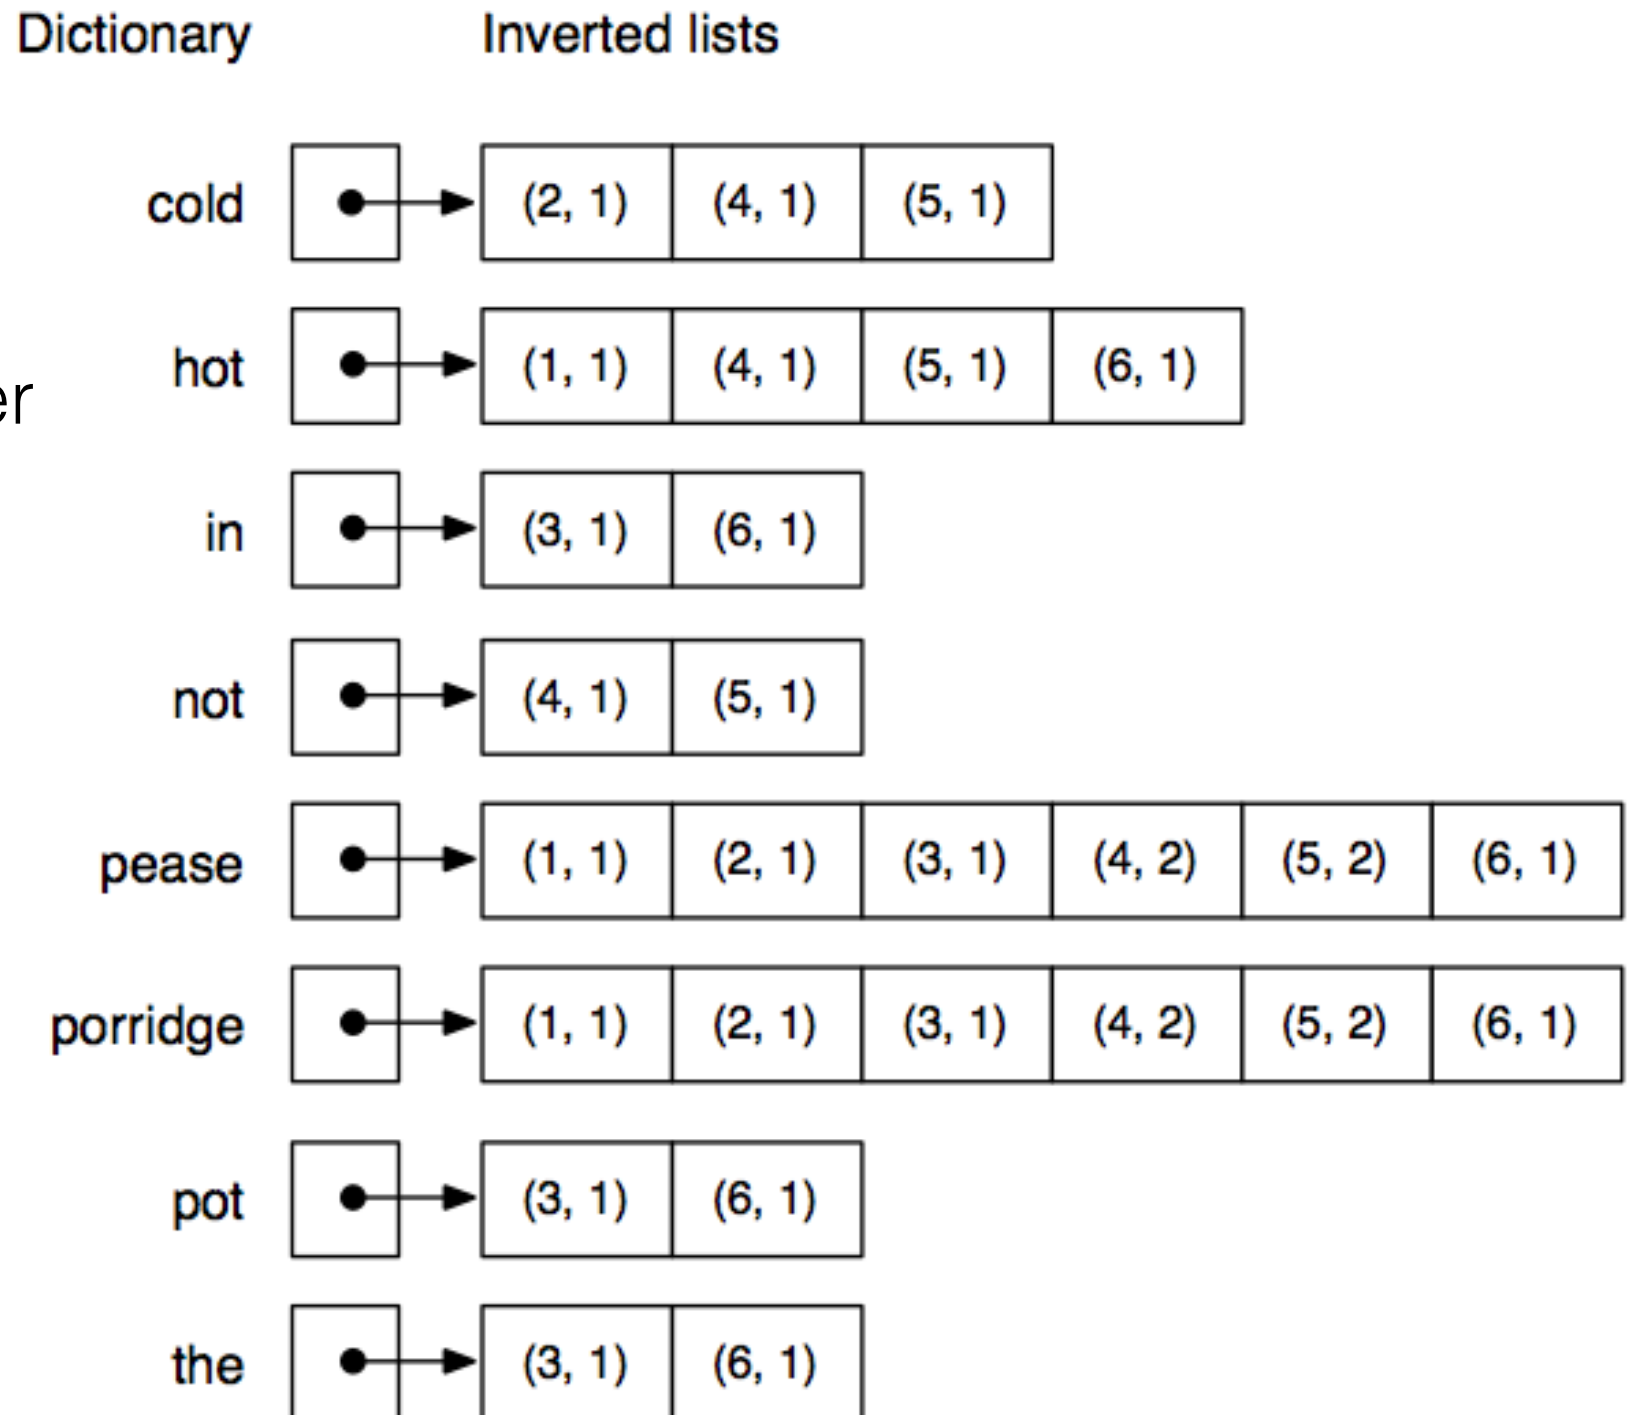
- original text: *Living in America*
- applying all: *liv america*
- in practice: *living in america*

Sample Document Collection

| Doc id | Text content |
|--------|---|
| 1 | pease porridge hot |
| 2 | pease porridge cold |
| 3 | pease porridge in the pot |
| 4 | pease porridge hot, pease porridge not cold |
| 5 | pease porridge cold, pease porridge not hot |
| 6 | pease porridge hot in the pot |


Inverted Index

- An inverted index is a representation for the document collection over which user queries are evaluated.
- It has two parts
 - a vocabulary index (dictionary)
 - inverted lists
 - document id
 - term information




Inverted Index

- Extensions
 - position lists: list of all positions a term occurs in a document



- skipping




- title, body, header, anchor text (inbound, outbound links)

Success Measure

- Quality measures
 - spam rate: fraction of spam pages in the index
 - duplicate rate: fraction of near duplicate web pages in the index
- Performance measures
 - compactness: size of the index in bytes
 - deployment cost: effort needed to create and deploy a new inverted index from scratch
 - update cost: time and space overhead of updating a document entry in the index

Compression

- Benefits
 - reduced space consumption
 - reduced transfer costs
 - increased posting list cache hit rate
 - Gap encoding
 - original: 17 18 28 40 44 47 56 58
 - gap encoded: 17 1 10 12 4 3 9 2
- 

Compression Algorithms

| Compression algorithm | Input sequence | Output | Parameters | Encoded values |
|-----------------------|---------------------|--------------|----------------|-------------------|
| Unary | gaps | bit-aligned | non-parametric | individual values |
| Gamma | gaps | bit-aligned | non-parametric | individual values |
| Delta | gaps | bit-aligned | non-parametric | individual values |
| Variable byte | gaps | byte-aligned | non-parametric | individual values |
| Golomb | gaps | bit-aligned | parametric | individual values |
| Simple-9 | gaps | word-aligned | parametric | blocks of values |
| PForDelta | gaps | bit-aligned | parametric | blocks of values |
| Binary interpolation | monotonic sequences | bit-aligned | parametric | bisections |
| Elias-Fano | monotonic sequences | bit-aligned | parametric | entire sequence |

Docid Reordering

- Goal: reassign document identifiers so that we obtain many small d-gaps, facilitating compression.

Id mapping:

1 → 1

2 → 9

3 → 2

4 → 7

5 → 8

6 → 3

7 → 5

8 → 6

9 → 4

Original lists:

L1: 1, 3, 6, 8, 9

L2: 2, 4, 5, 6, 9

L3: 3, 6, 7, 9

Original d-gaps:

L1: 2, 3, 2, 1

L2: 2, 1, 1, 3

L3: 3, 1, 2

Reordered lists:

L1: 1, 2, 3, 4, 6

L2: 3, 4, 7, 8, 9

L3: 2, 3, 4, 5

New d-gaps:

L1: 1, 1, 1, 2

L2: 1, 3, 1, 1

L3: 1, 1, 1

Docid Reordering

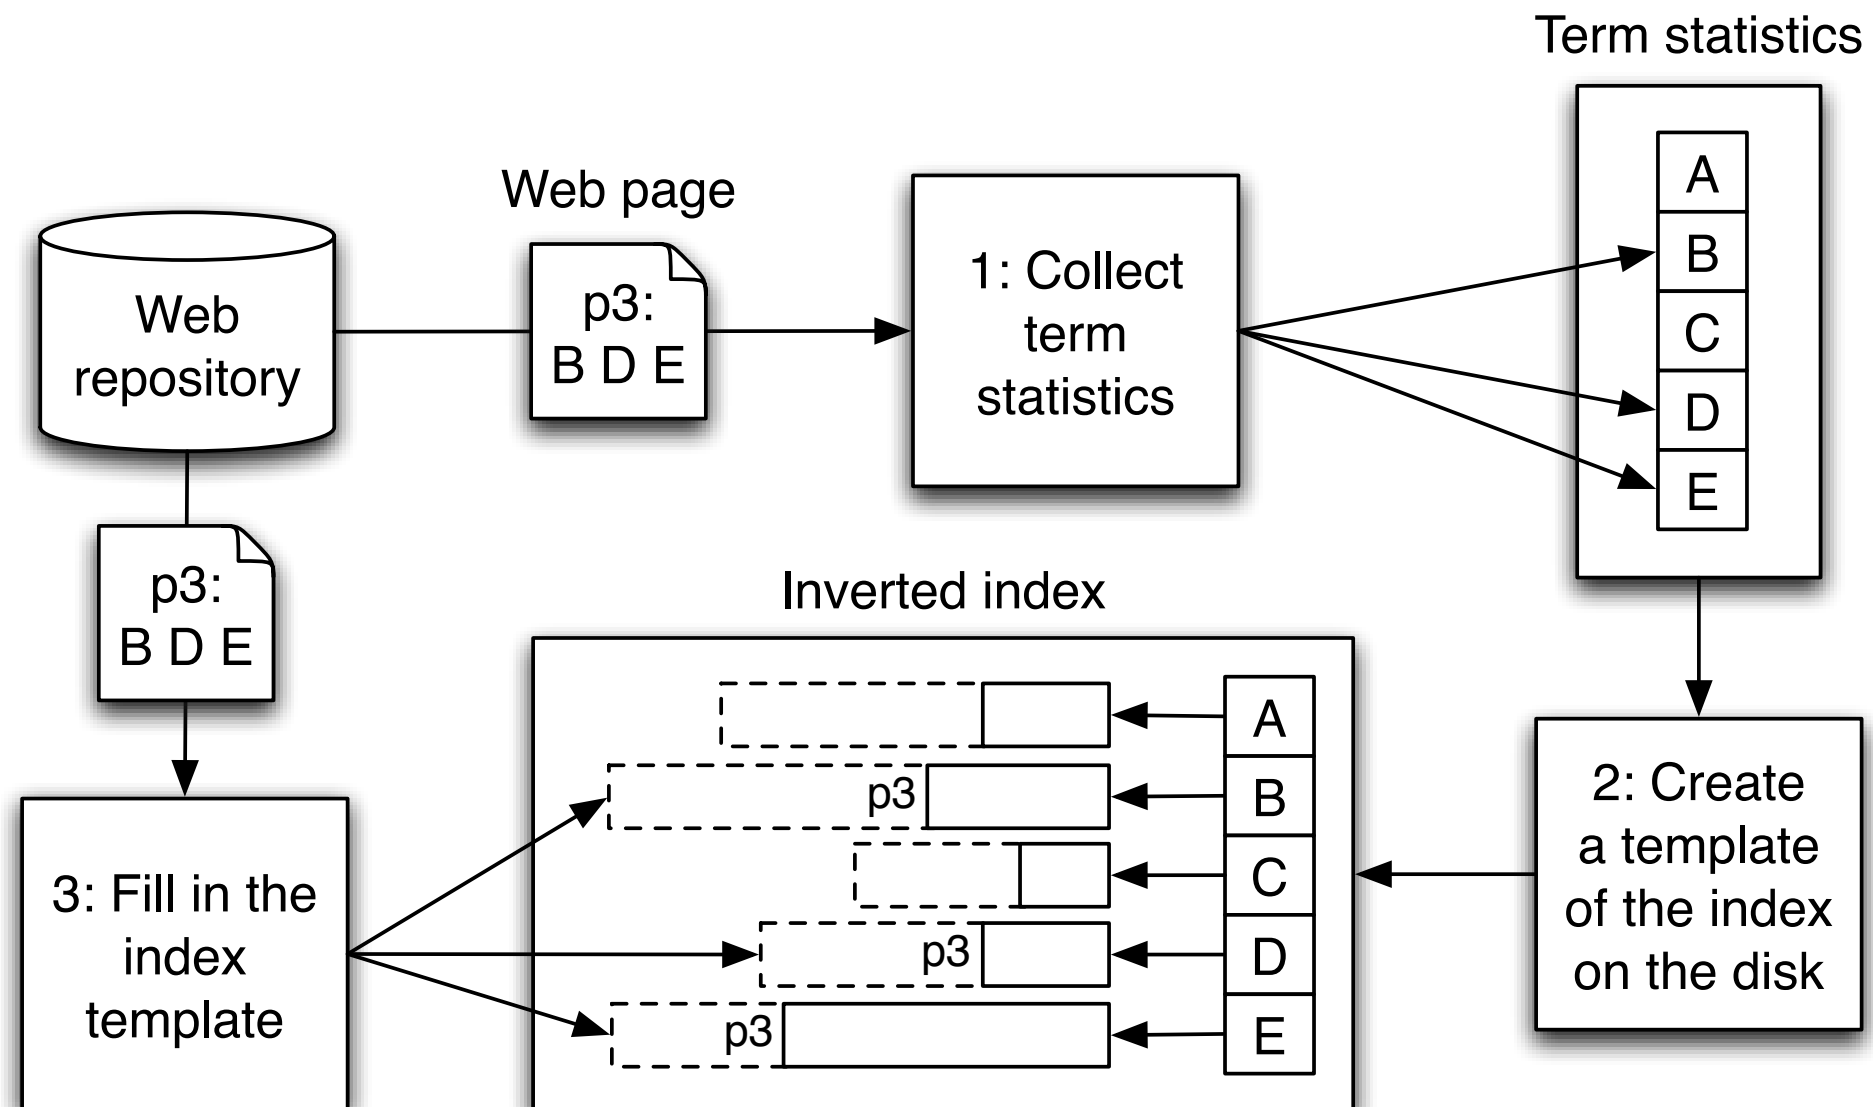
- Techniques
 - traversal of document similarity graph
 - formulated as the traveling salesman problem
 - clustering similar documents
 - assigns nearby ids to documents in the same cluster
 - sorting URLs alphabetically and assigning ids in that order
 - idea: pages from the same site have high textual overlap
 - simple yet effective
 - only applicable to web page collections

Index Construction

- Equivalent to computing the transpose of a matrix.
- In-memory techniques do not work well with web-scale data.
- Techniques
 - two-phase
 - one-phase

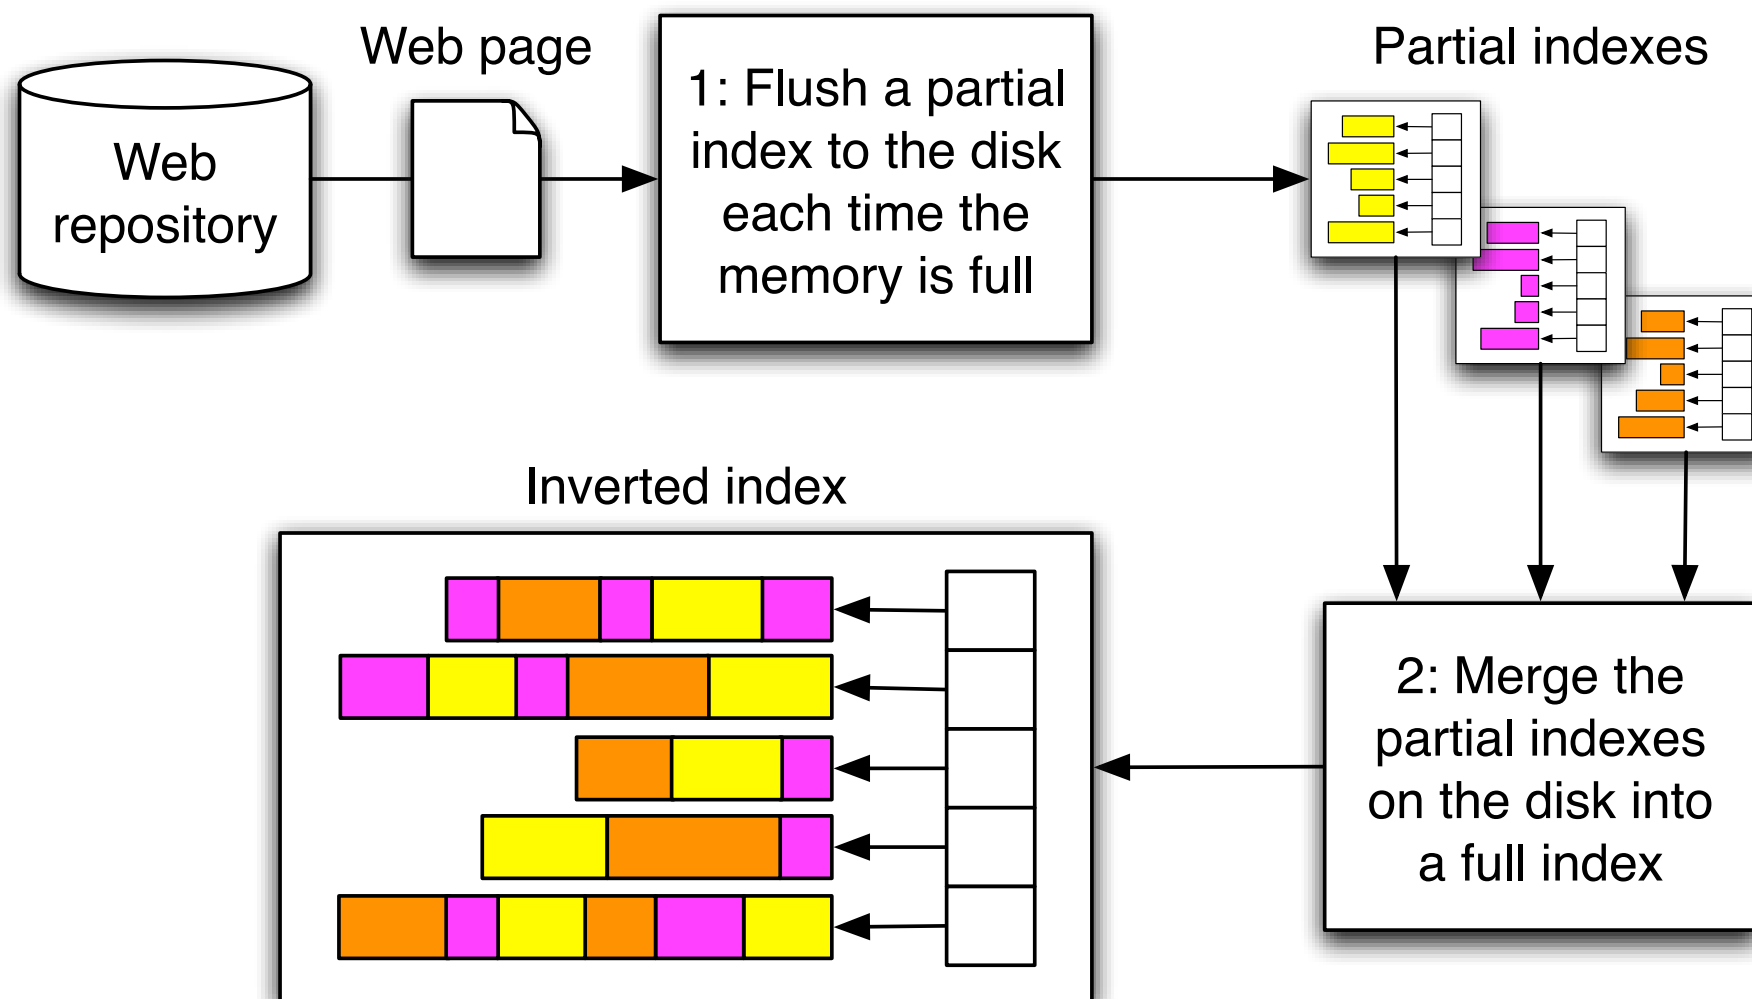
Two Phase

- First phase: read the collection and allocate a skeleton for the index.
- Second phase: fill the posting lists.



One Phase

- Keep reading documents and building an in-memory index.
- Each time the memory is full, flush the index to the disk.
- Merge all on-disk indexes into a single index in a final step.




Index Maintenance

- Grow a new (delta) index in the memory; each time the memory is full, flush the in-memory index to disk.
- Techniques
 - no merge
 - incremental update
 - immediate merge
 - lazy merge


No Merge

- Flushed index is written to disk as a separate index.
- Increases fragmentation and query processing time.
- Eventually requires merging all on-disk indexes or rebuilding.




Incremental Update

- Each inverted list contains additional empty space at the end.
- New documents are appended to the empty space in the list.
- If the extra space allocated in an inverted list is full.
 - inverted list may be reallocated on disk
 - inverted list is maintained in multiple fragments on disk




Immediate Merge

- Delta index is immediately merged to the old index and written to a new location on disk.
- Only one copy of the index is maintained on disk.




Lazy Merge

- Maintains multiple generations of the index on disk.
- Index generations are lazily merged.



Inverted Index Partitioning/Replication

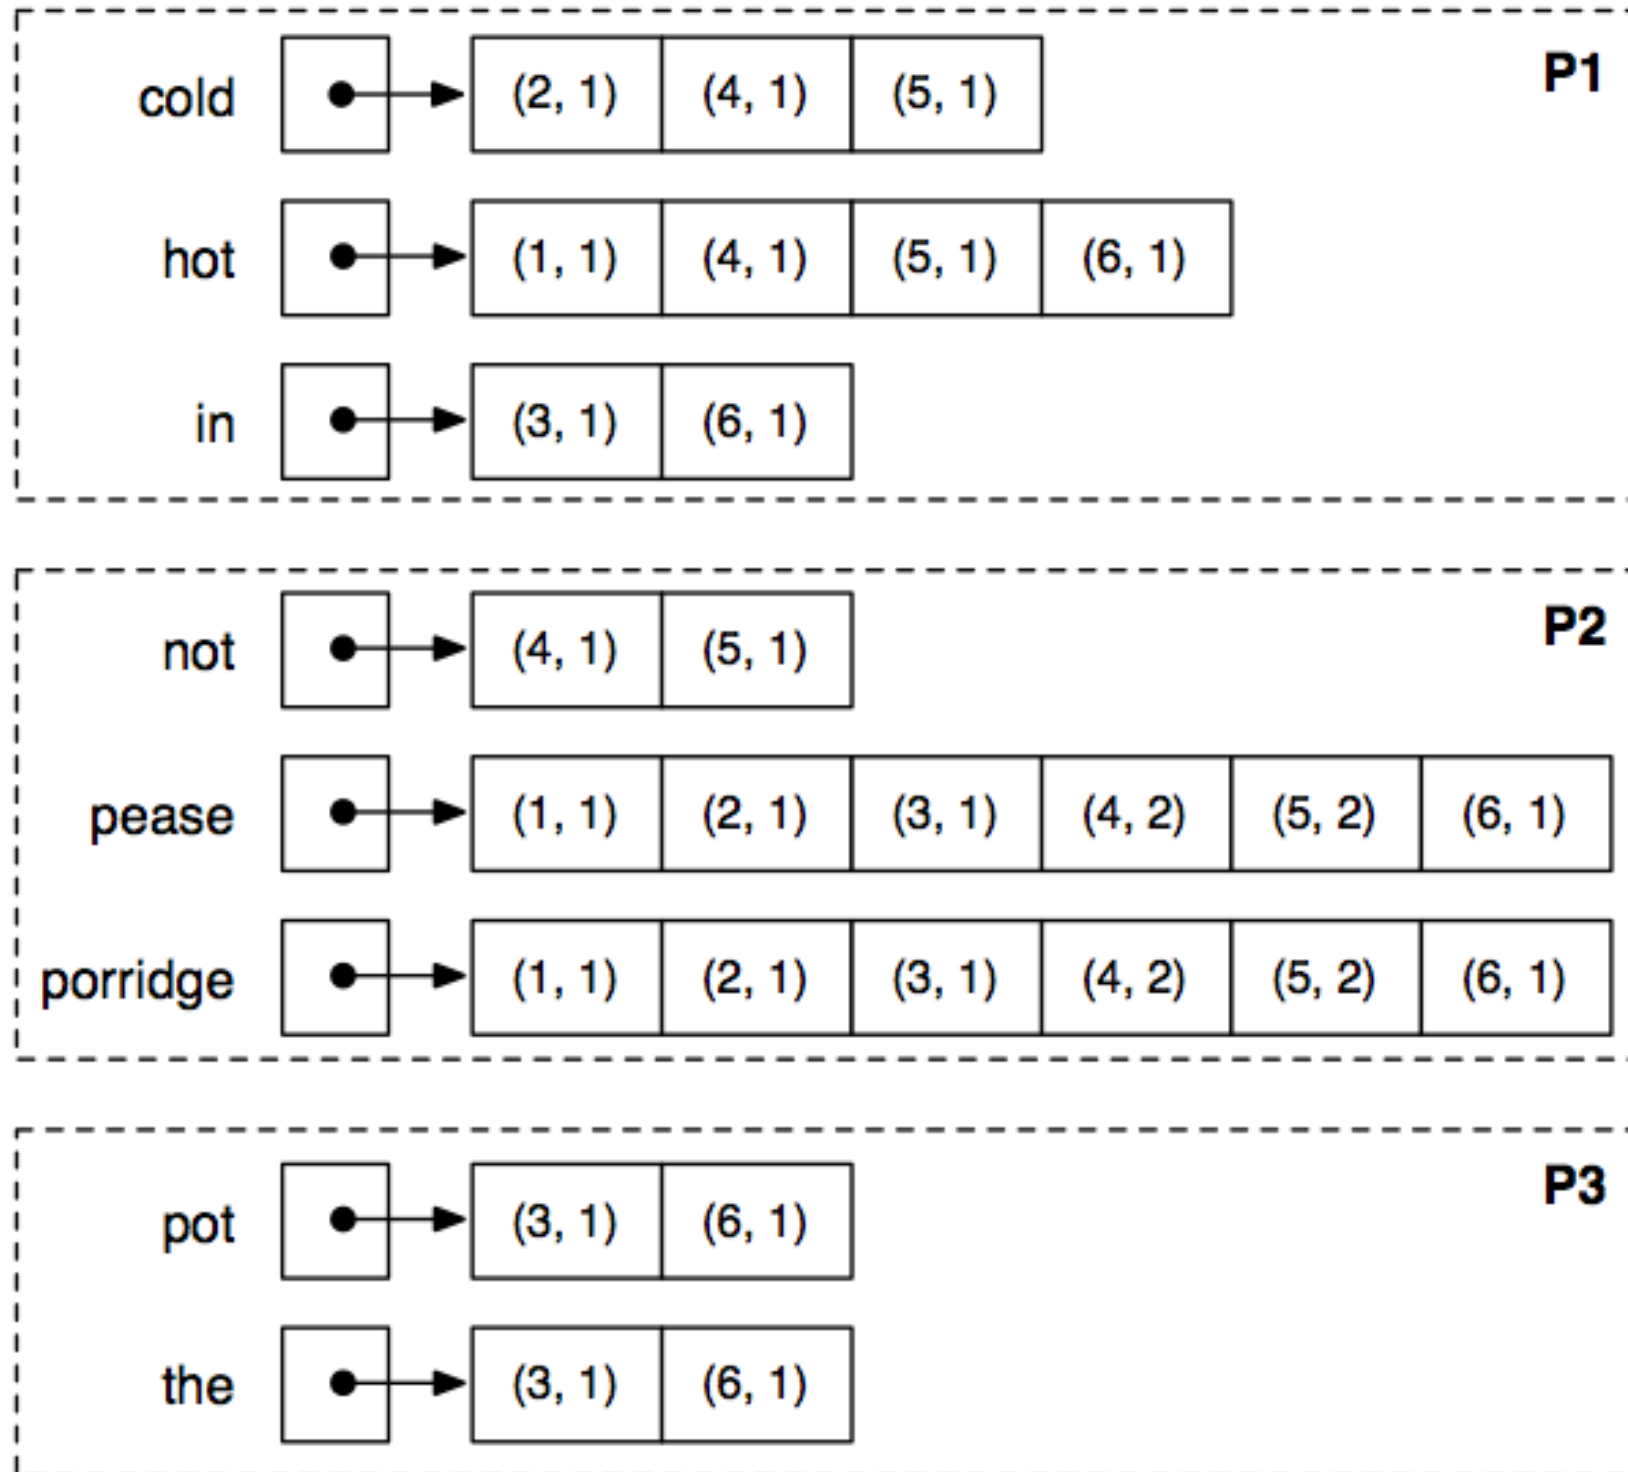
- In practice, the inverted index is
 - partitioned on thousands of computers in a large search cluster
 - reduces query response times
 - allows scaling with increasing collection size
 - replicated on tens of search clusters
 - increases query processing throughput
 - allows scaling with increasing query volume
 - provides fault tolerance



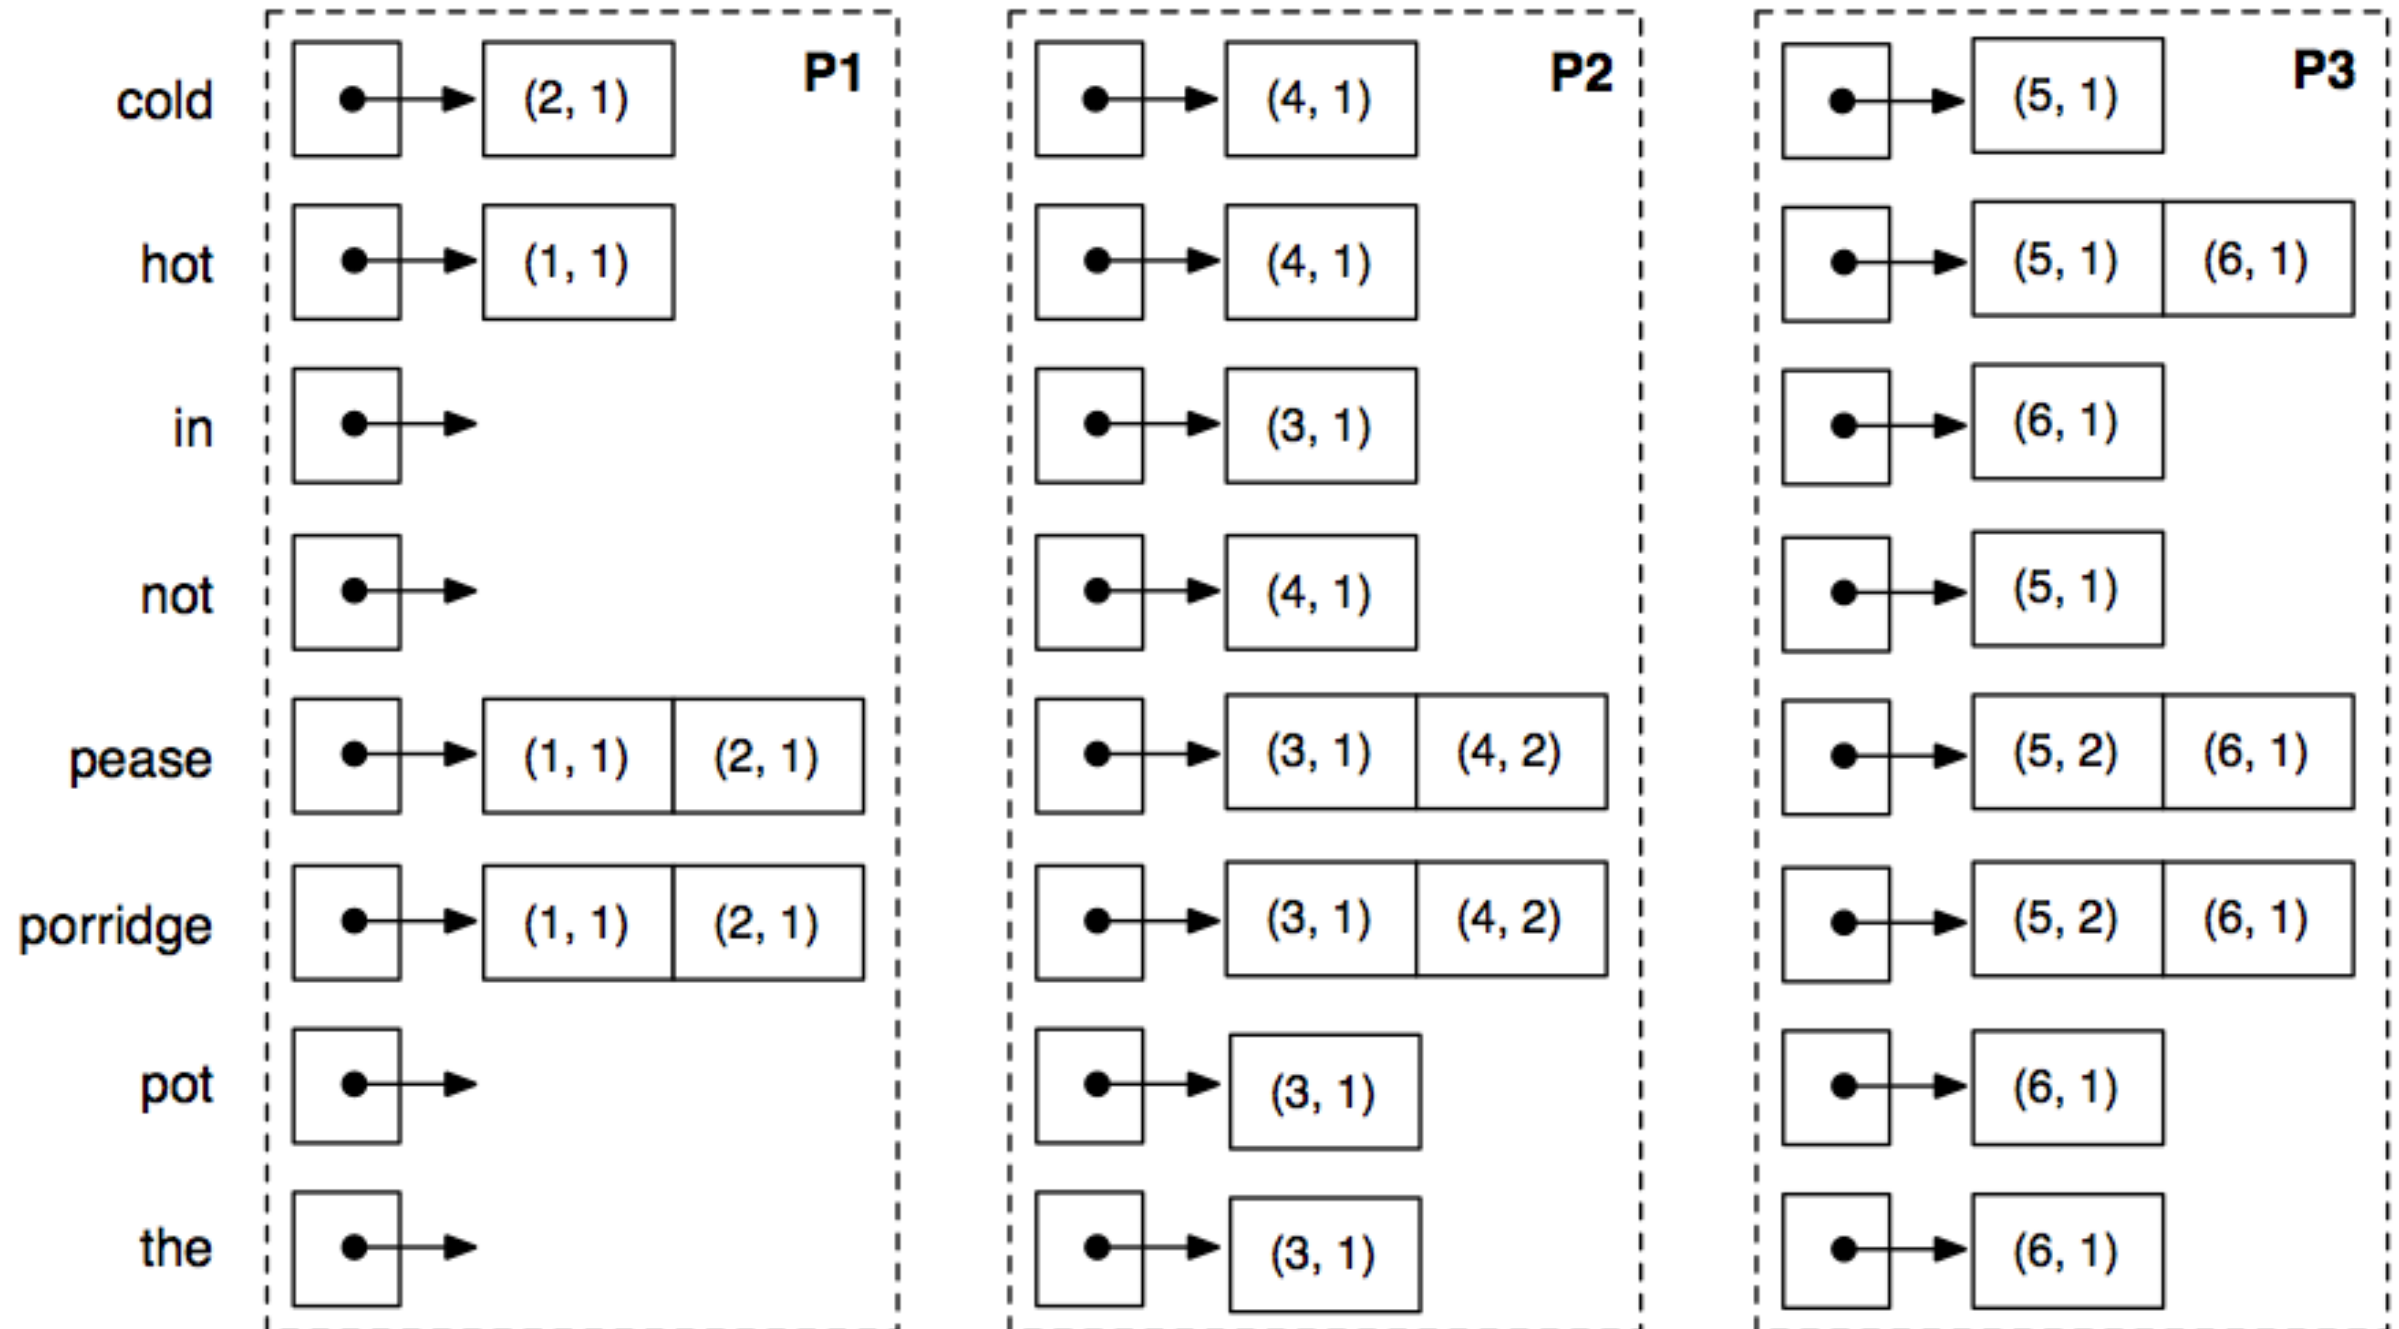
Inverted Index Partitioning

- Two alternatives for partitioning an inverted index
 - term-based partitioning
 - T inverted lists are distributed across P processors
 - each processor is responsible for processing the postings of a mutually disjoint subset of inverted lists assigned to itself
 - single disk access per query term
 - document-based partitioning
 - N documents are distributed across P processors
 - each processor is responsible for processing the postings of a mutually disjoint subset of documents assigned to itself
 - multiple (parallel) disk accesses per query term

Term-Based Index Partitioning



Document-Based Index Partitioning



Comparison of Index Partitioning Approaches

| | Document-based | Term-based |
|------------------------------|----------------|------------|
| Space consumption | Higher | Lower |
| Number of disk accesses | Higher | Lower |
| Concurrency | Lower | Higher |
| Computational load imbalance | Lower | Higher |
| Max. posting list I/O time | Lower | Higher |
| Cost of index building | Lower | Higher |
| Maintenance cost | Lower | Higher |

Inverted Index Partitioning

- In practice, document-based partitioning is used
 - easier to build and maintain
 - low inter-query-processing concurrency, but good load balance
 - low query processing time
 - high throughput is achieved by replication
 - more fault tolerant
- Hybrid techniques are possible (e.g., term partitioning inside a document sub-collection).

Indexing with MapReduce

- **Map over documents**
 - Emit *term* as key, (*docno*, *tf*) posting as value
 - Emit other information as necessary (e.g., term position)
- **Group postings by term**
 - Sort the postings (by docid)
 - Write postings to disk

Indexing with MapReduce (I)

```

1: class MAPPER
2:   method MAP(docid  $n$ , doc  $d$ )
3:      $H \leftarrow$  new ASSOCIATIVEARRAY                                ▷ histogram to hold term frequencies
4:     for all term  $t \in$  doc  $d$  do                                     ▷ processes the doc, e.g., tokenization and stopword removal
5:        $H\{t\} \leftarrow H\{t\} + 1$ 
6:     for all term  $t \in H$  do
7:       EMIT(term  $t$ , posting  $\langle n, H\{t\} \rangle$ )                       ▷ emits individual postings

1: class REDUCER
2:   method REDUCE(term  $t$ , postings  $[\langle n_1, f_1 \rangle \dots]$ )
3:      $P \leftarrow$  new LIST
4:     for all  $\langle n, f \rangle \in$  postings  $[\langle n_1, f_1 \rangle \dots]$  do
5:        $P$ .APPEND( $\langle n, f \rangle$ )                                         ▷ appends postings unsorted
6:        $P$ .SORT()                                                    ▷ sorts for compression
7:       EMIT(term  $t$ , postingsList  $P$ )

```

Indexing with MapReduce (II)

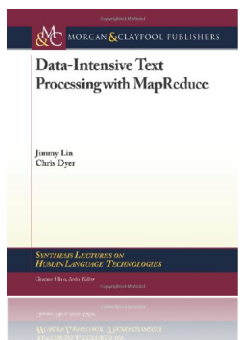
```

1: class MAPPER
2:   method MAP(docid  $n$ , doc  $d$ )
3:      $H \leftarrow$  new ASSOCIATIVEARRAY
4:     for all term  $t \in$  doc  $d$  do                                     ▷ builds a histogram of term frequencies
5:        $H\{t\} \leftarrow H\{t\} + 1$ 
6:     for all term  $t \in H$  do
7:       EMIT(tuple  $\langle t, n \rangle$ , tf  $H\{t\}$ )                             ▷ emits individual postings, with a tuple as the key

1: class PARTITIONER
2:   method PARTITION(tuple  $\langle t, n \rangle$ , tf  $f$ )
3:     return HASH( $t$ ) mod NumOfReducers                             ▷ keys of same term are sent to same reducer

1: class REDUCER
2:   method INITIALIZE
3:      $t_{prev} \leftarrow \emptyset$ 
4:      $P \leftarrow$  new POSTINGSLIST
5:   method REDUCE(tuple  $\langle t, n \rangle$ , tf  $[f]$ )
6:     if  $t \neq t_{prev} \wedge t_{prev} \neq \emptyset$  then
7:       EMIT(term  $t$ , postings  $P$ )                                     ▷ emits postings list of term  $t_{prev}$ 
8:        $P$ .RESET()
9:        $P$ .APPEND( $\langle n, f \rangle$ )                                         ▷ appends postings in sorted order
10:       $t_{prev} \leftarrow t$ 
11:   method CLOSE
12:     EMIT(term  $t$ , postings  $P$ )                                     ▷ emits last postings list from this reducer

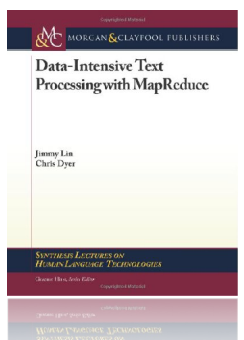
```



Indexing with MapReduce (III)

```

1: class MAPPER
2:   method INITIALIZE
3:      $M \leftarrow \text{new ASSOCIATIVEARRAY}$                                 ▷ holds partial lists of postings
4:   method MAP(docid  $n$ , doc  $d$ )
5:      $H \leftarrow \text{new ASSOCIATIVEARRAY}$                                 ▷ builds a histogram of term frequencies
6:     for all term  $t \in \text{doc } d$  do
7:        $H\{t\} \leftarrow H\{t\} + 1$ 
8:     for all term  $t \in H$  do
9:        $M\{t\}.\text{ADD}(\text{posting } \langle n, H\{t\} \rangle)$                         ▷ adds a posting to partial postings lists
10:    if MEMORYFULL() then
11:      FLUSH()
12:  method FLUSH                                ▷ flushes partial lists of postings as intermediate output
13:    for all term  $t \in M$  do
14:       $P \leftarrow \text{SORTANDENCODEPOSTINGS}(M\{t\})$ 
15:       $\text{EMIT}(\text{term } t, \text{postingsList } P)$ 
16:     $M.\text{CLEAR}()$ 
17:  method CLOSE
18:    FLUSH()
  
```



Indexing with MapReduce (III)

```

1: class REDUCER
2:   method REDUCE(term  $t$ , postingsLists  $[P_1, P_2, \dots]$ )
3:      $P_f \leftarrow$  new LIST                                ▷ temporarily stores partial lists of postings
4:      $R \leftarrow$  new LIST                                  ▷ stores merged partial lists of postings
5:     for all  $P \in$  postingsLists  $[P_1, P_2, \dots]$  do
6:        $P_f.$ ADD( $P$ )
7:       if MEMORYNEARLYFULL() then
8:          $R.$ ADD(MERGELists( $P_f$ ))
9:          $P_f.$ CLEAR()
10:     $R.$ ADD(MERGELists( $P_f$ ))
11:    EMIT(term  $t$ , postingsList MERGELists( $R$ ))           ▷ emits fully merged postings list of term  $t$ 

```

Pagerank

- **Random Surfers**

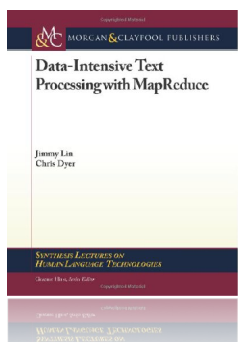
- User starts at a random Web page
- User randomly clicks on links, surfing from page to page

- **Pagerank**

- Characterizes the amount of time spent on any given page
- Mathematically, a probability distribution over pages

- **Web Ranking**



- One of thousands of features used in web search



Definition

- Given page x with inlinks t_1, \dots, t_n , where
 - $C(t)$ is the out-degree of link t
 - α is probability of random jump
 - N is the total number of nodes in the graph

$$PR(x) = \alpha \left(\frac{1}{N} \right) + (1 - \alpha) \sum_{i=1}^n \frac{PR(t_i)}{C(t_i)}$$


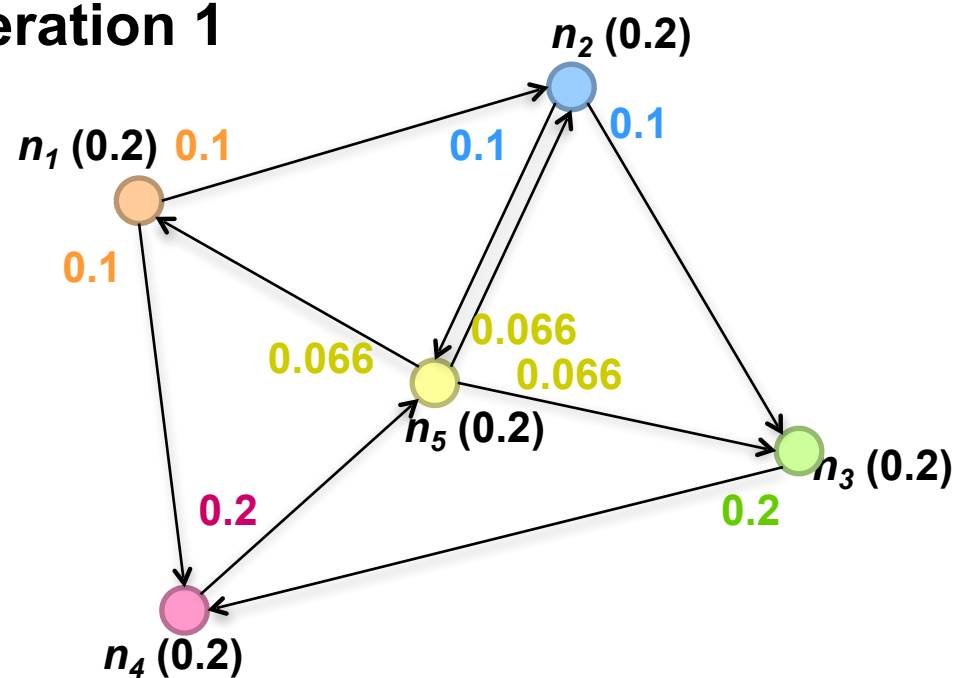


Algorithm Sketch

- Start with seed PR_i values
- Each page distributes PR_i mass to all pages it links to
- Each target page adds up mass from in-bound links to compute $PR_i + 1$
- Iterate until values converge


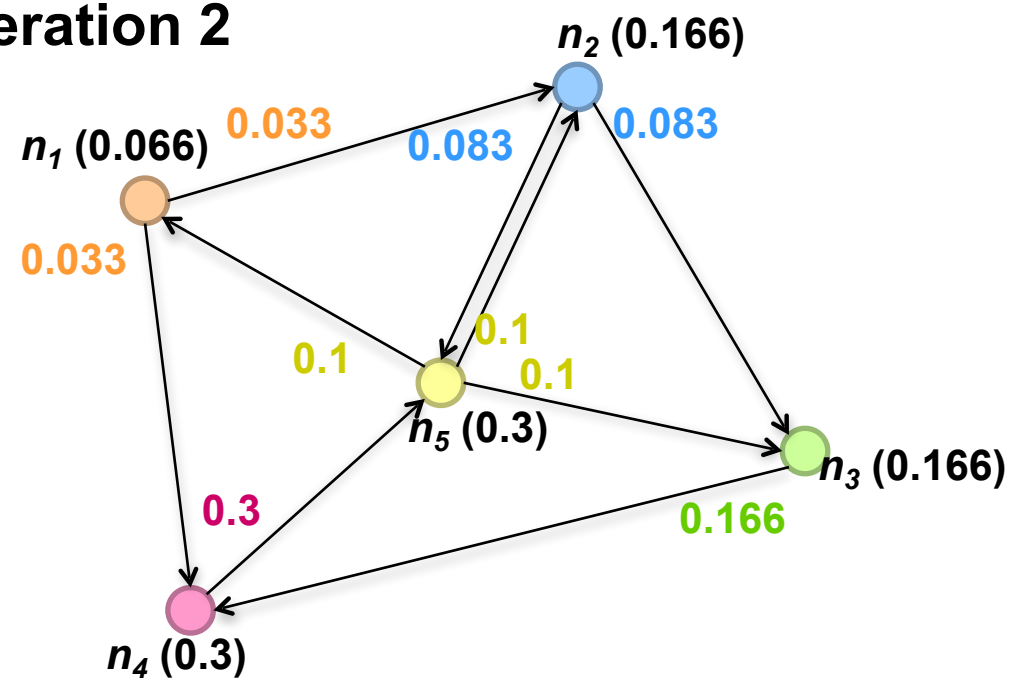
Simplified Algorithm Example (I)

Iteration 1




Simplified Algorithm Example (II)

Iteration 2




Pagerank in MapReduce (I)

Map



Reduce



Pagerank in MapReduce (II)

```

1: class MAPPER
2:   method MAP(nid  $n$ , node  $N$ )
3:      $p \leftarrow N.PAGERANK / |N.ADJACENCYLIST|$ 
4:     EMIT(nid  $n$ ,  $N$ )                                ▷ Pass along graph structure
5:     for all nodeid  $m \in N.ADJACENCYLIST$  do
6:       EMIT(nid  $m$ ,  $p$ )                                ▷ Pass PageRank mass to neighbors

1: class REDUCER
2:   method REDUCE(nid  $m$ , [ $p_1, p_2, \dots$ ])
3:      $M \leftarrow \emptyset$ 
4:     for all  $p \in$  counts [ $p_1, p_2, \dots$ ] do
5:       if ISNODE( $p$ ) then
6:          $M \leftarrow p$                                 ▷ Recover graph structure
7:       else
8:          $s \leftarrow s + p$                                 ▷ Sums incoming PageRank contributions
9:      $M.PAGERANK \leftarrow s$ 
10:    EMIT(nid  $m$ , node  $M$ )

```