

Disclaimer

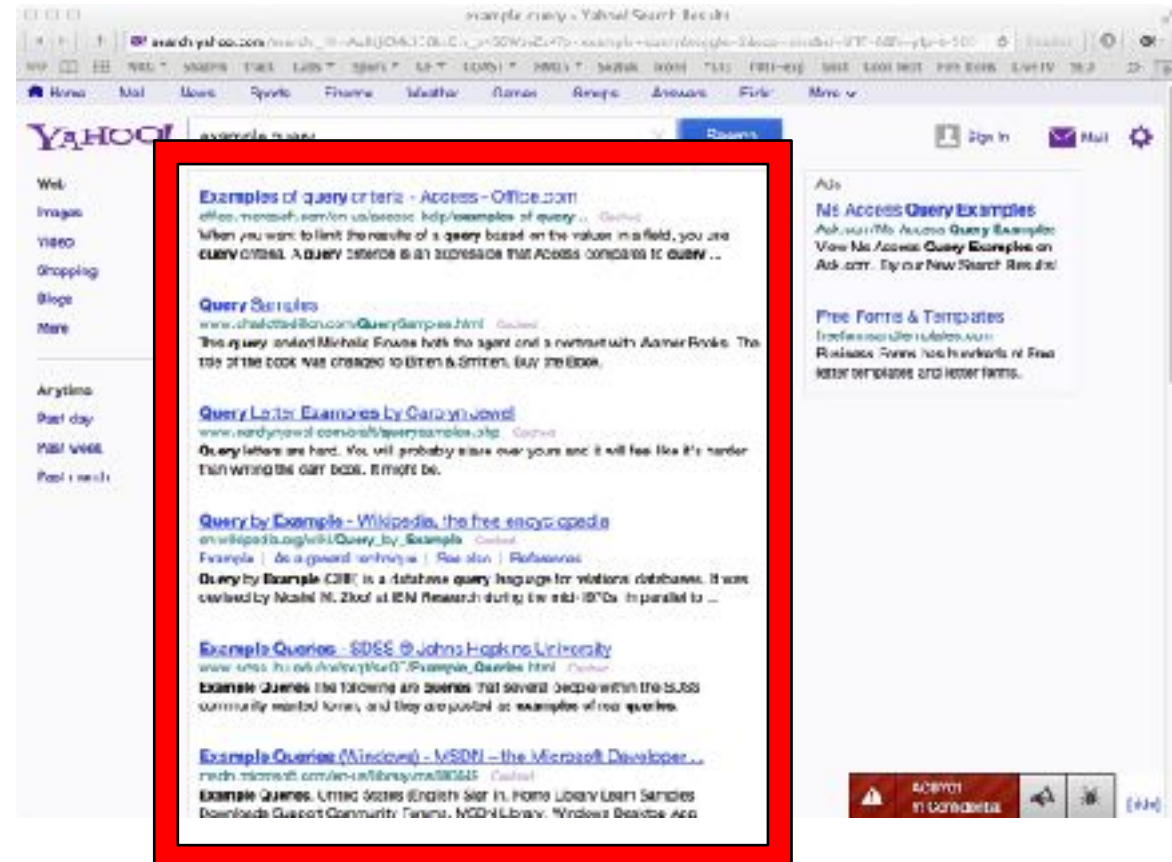
- Some of these slides have been prepared by B. Barla Cambazoglu, Director of Applied Science, NTENT Inc.
- Some of these slides have been prepared by Craig Macdonald (University of Glasgow) and Nicola Tonellotto (ISTI-CNR)
- Used with permission, please do not redistribute in any form.
- Please refer to:



Query Processing

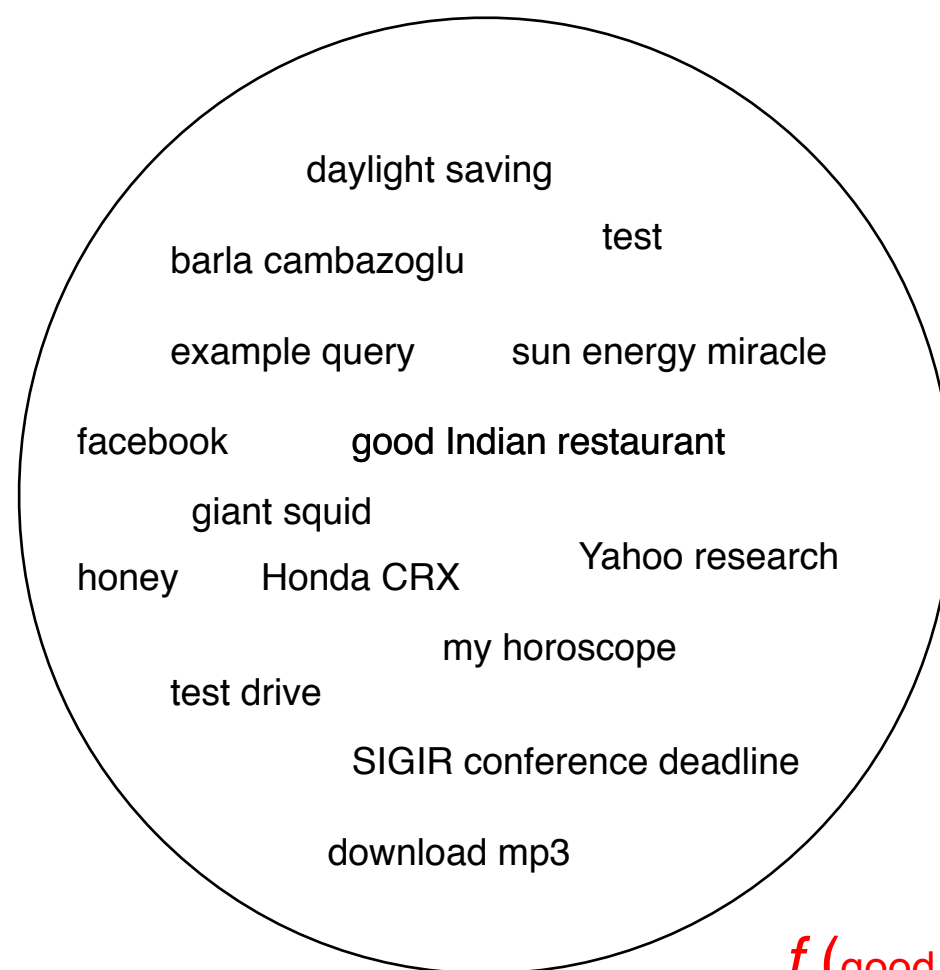
- Query processing is the problem of generating the best-matching answers (typically, top 10 documents) to a given user query, spending the least amount of time.

- Our focus: creating 10 blue links as an answer to a user query.



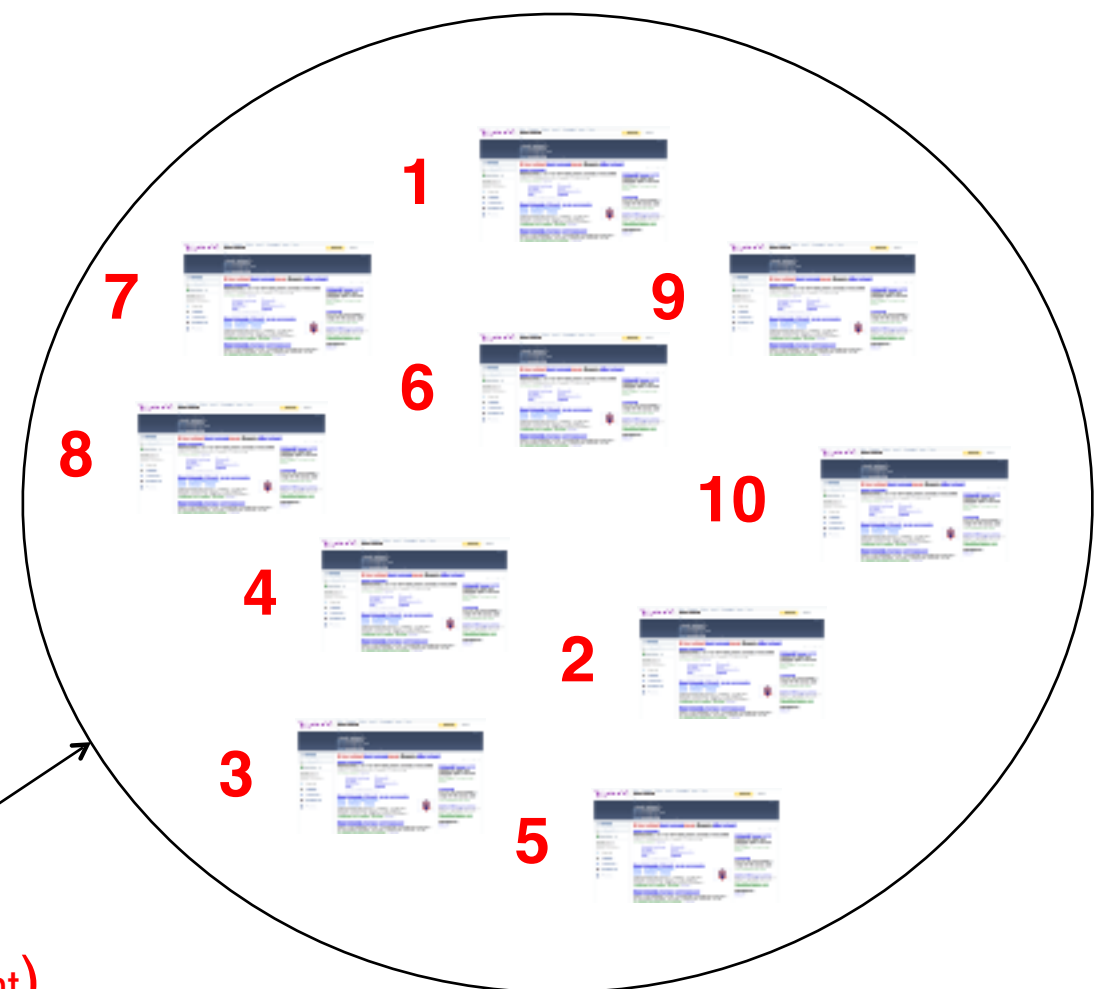
Web Search

- Web search is a sorting problem!



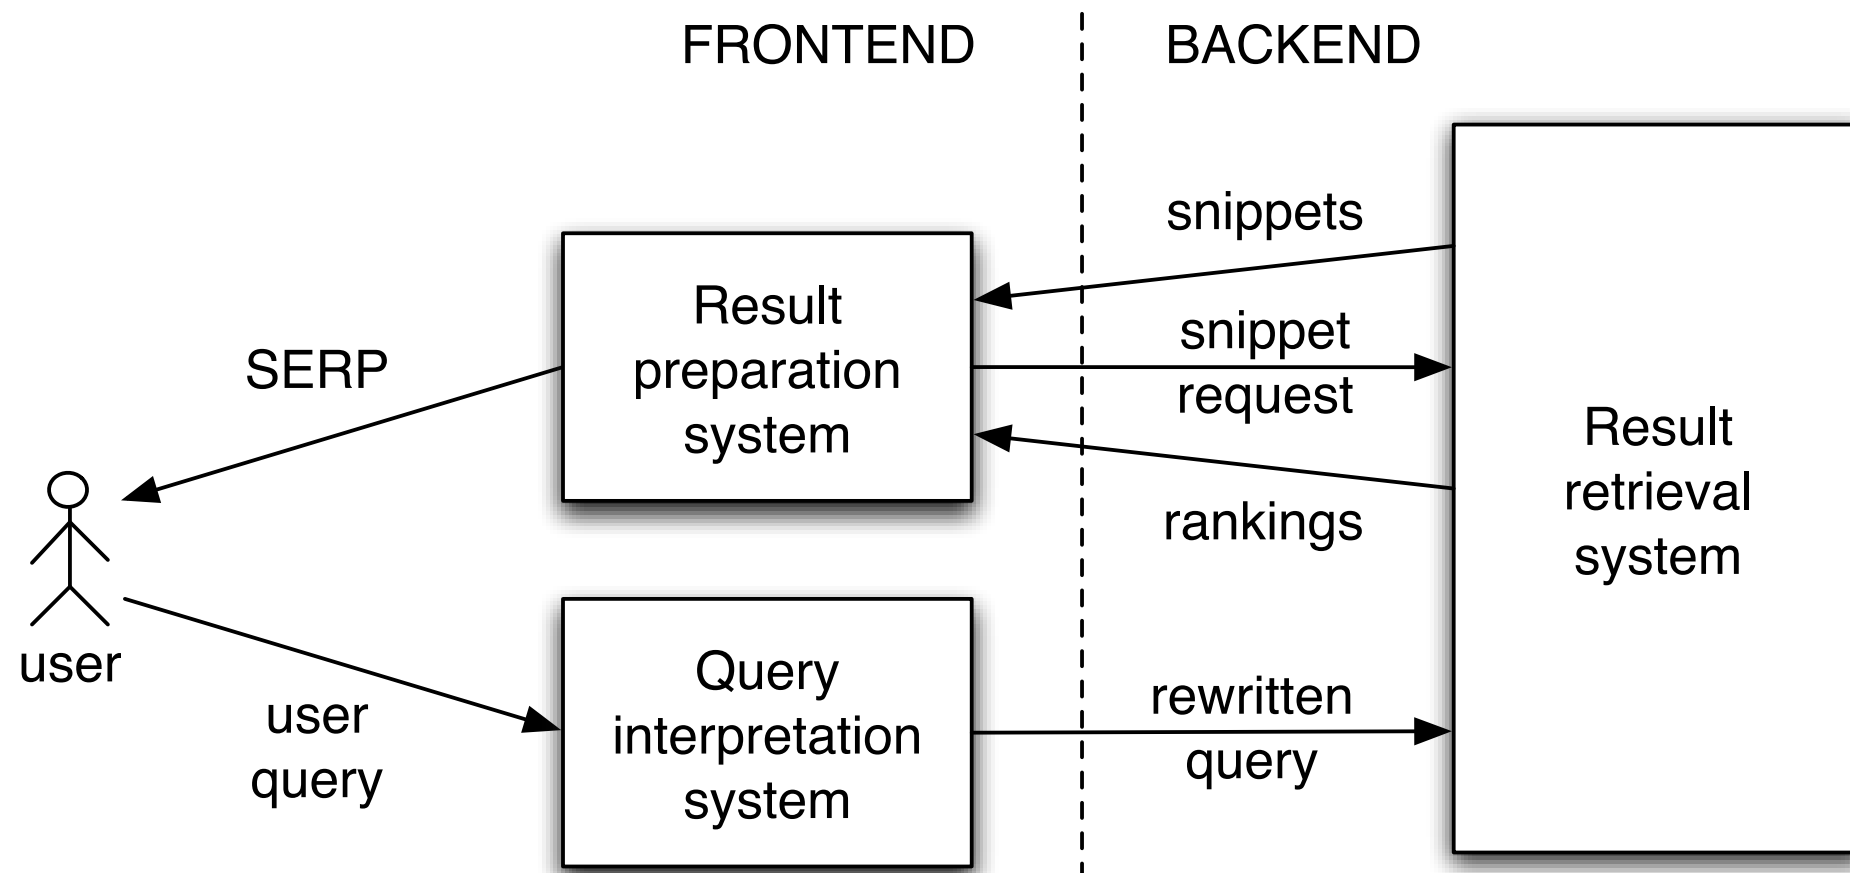
user queries

f (good Indian restaurant)

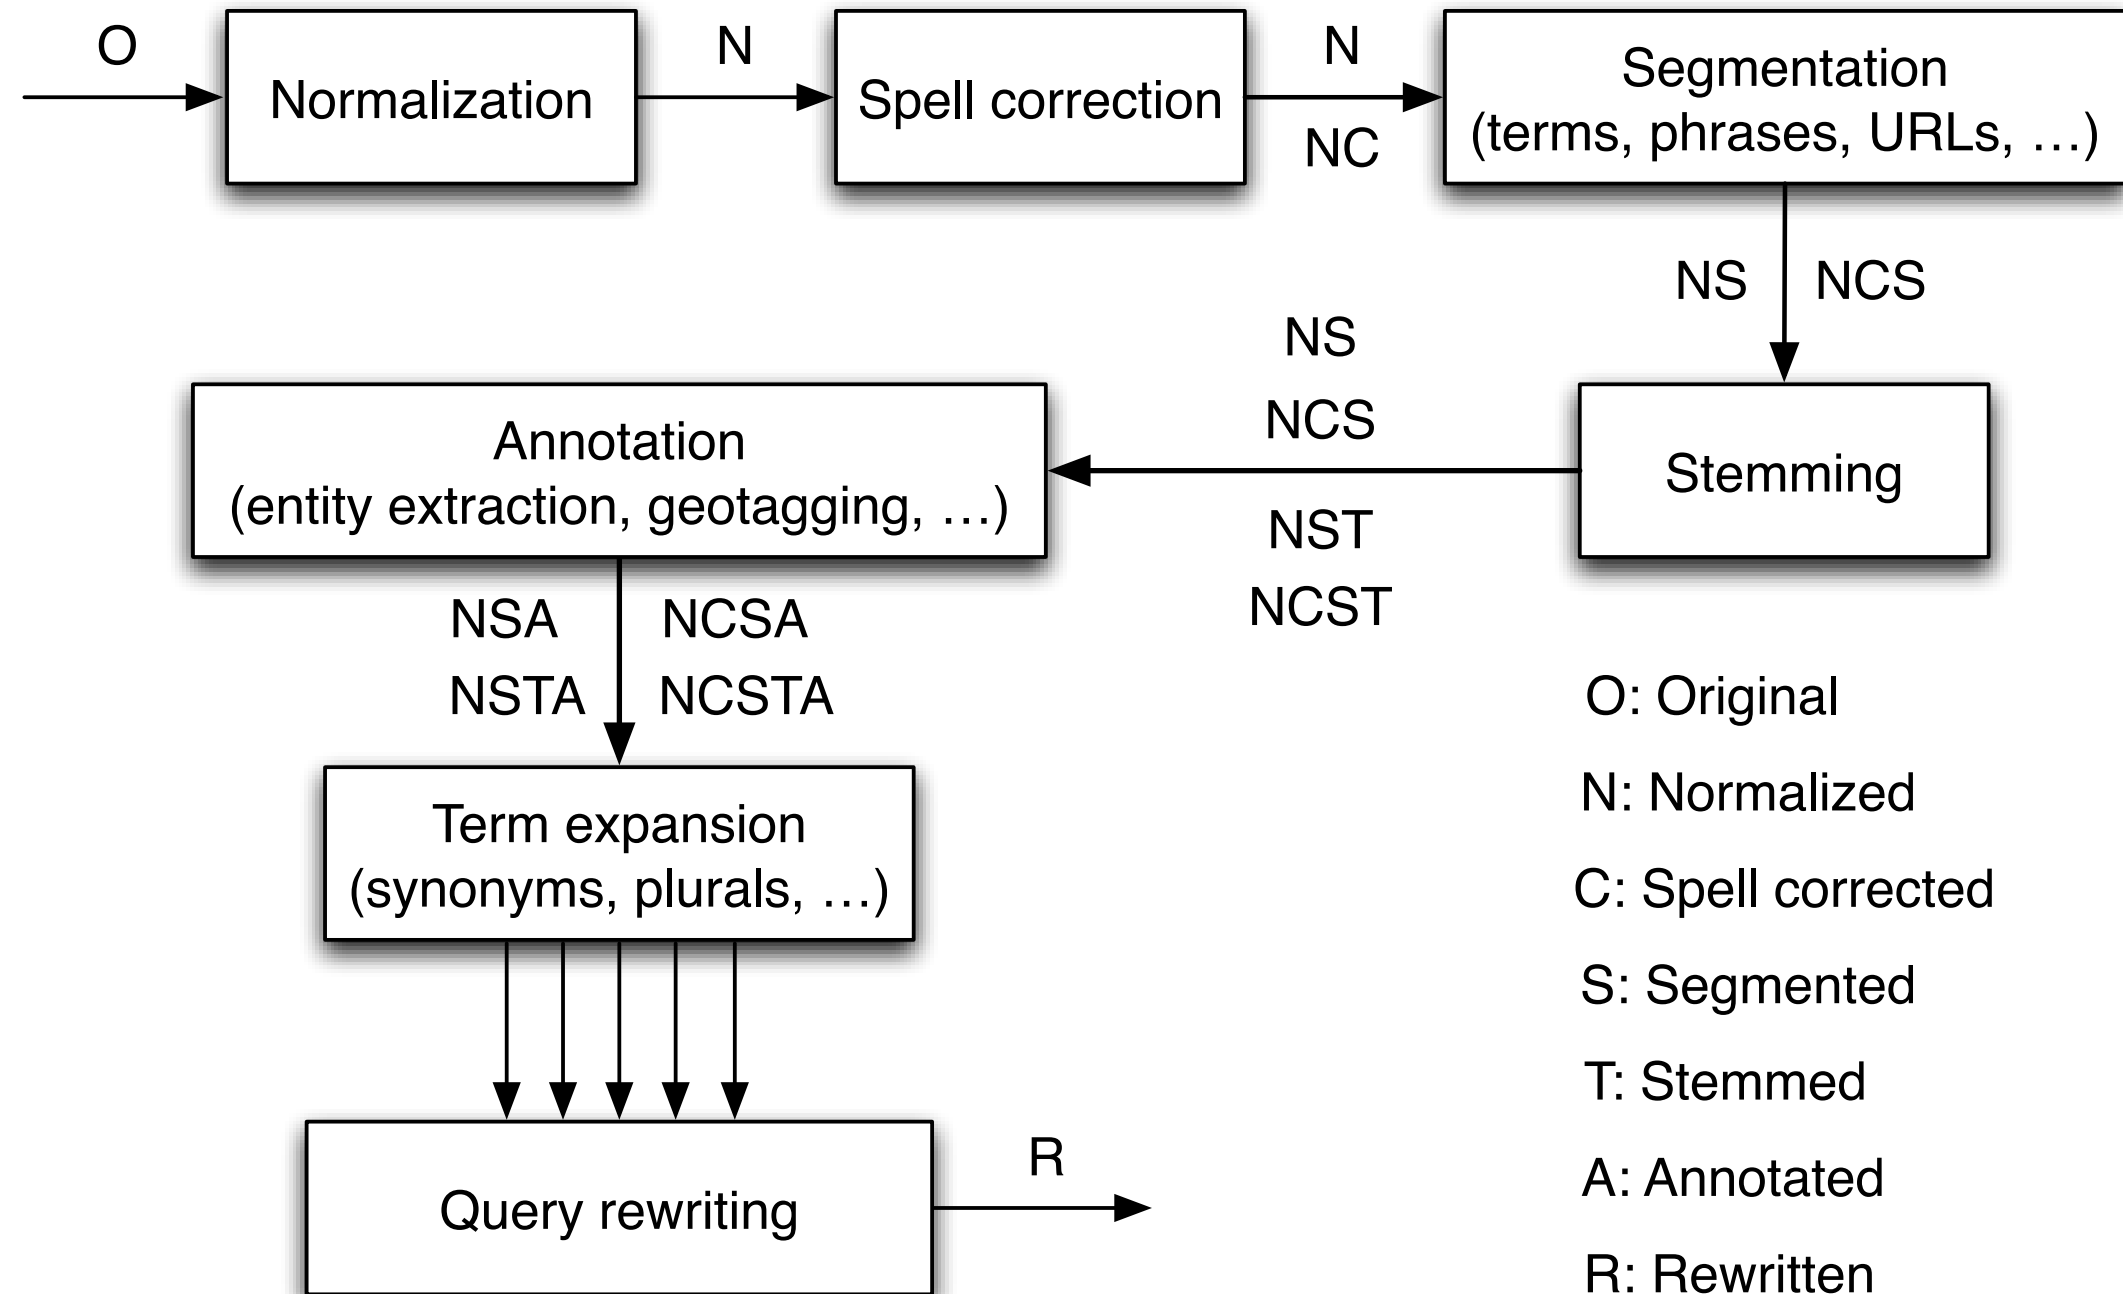


the Web

Query Processing



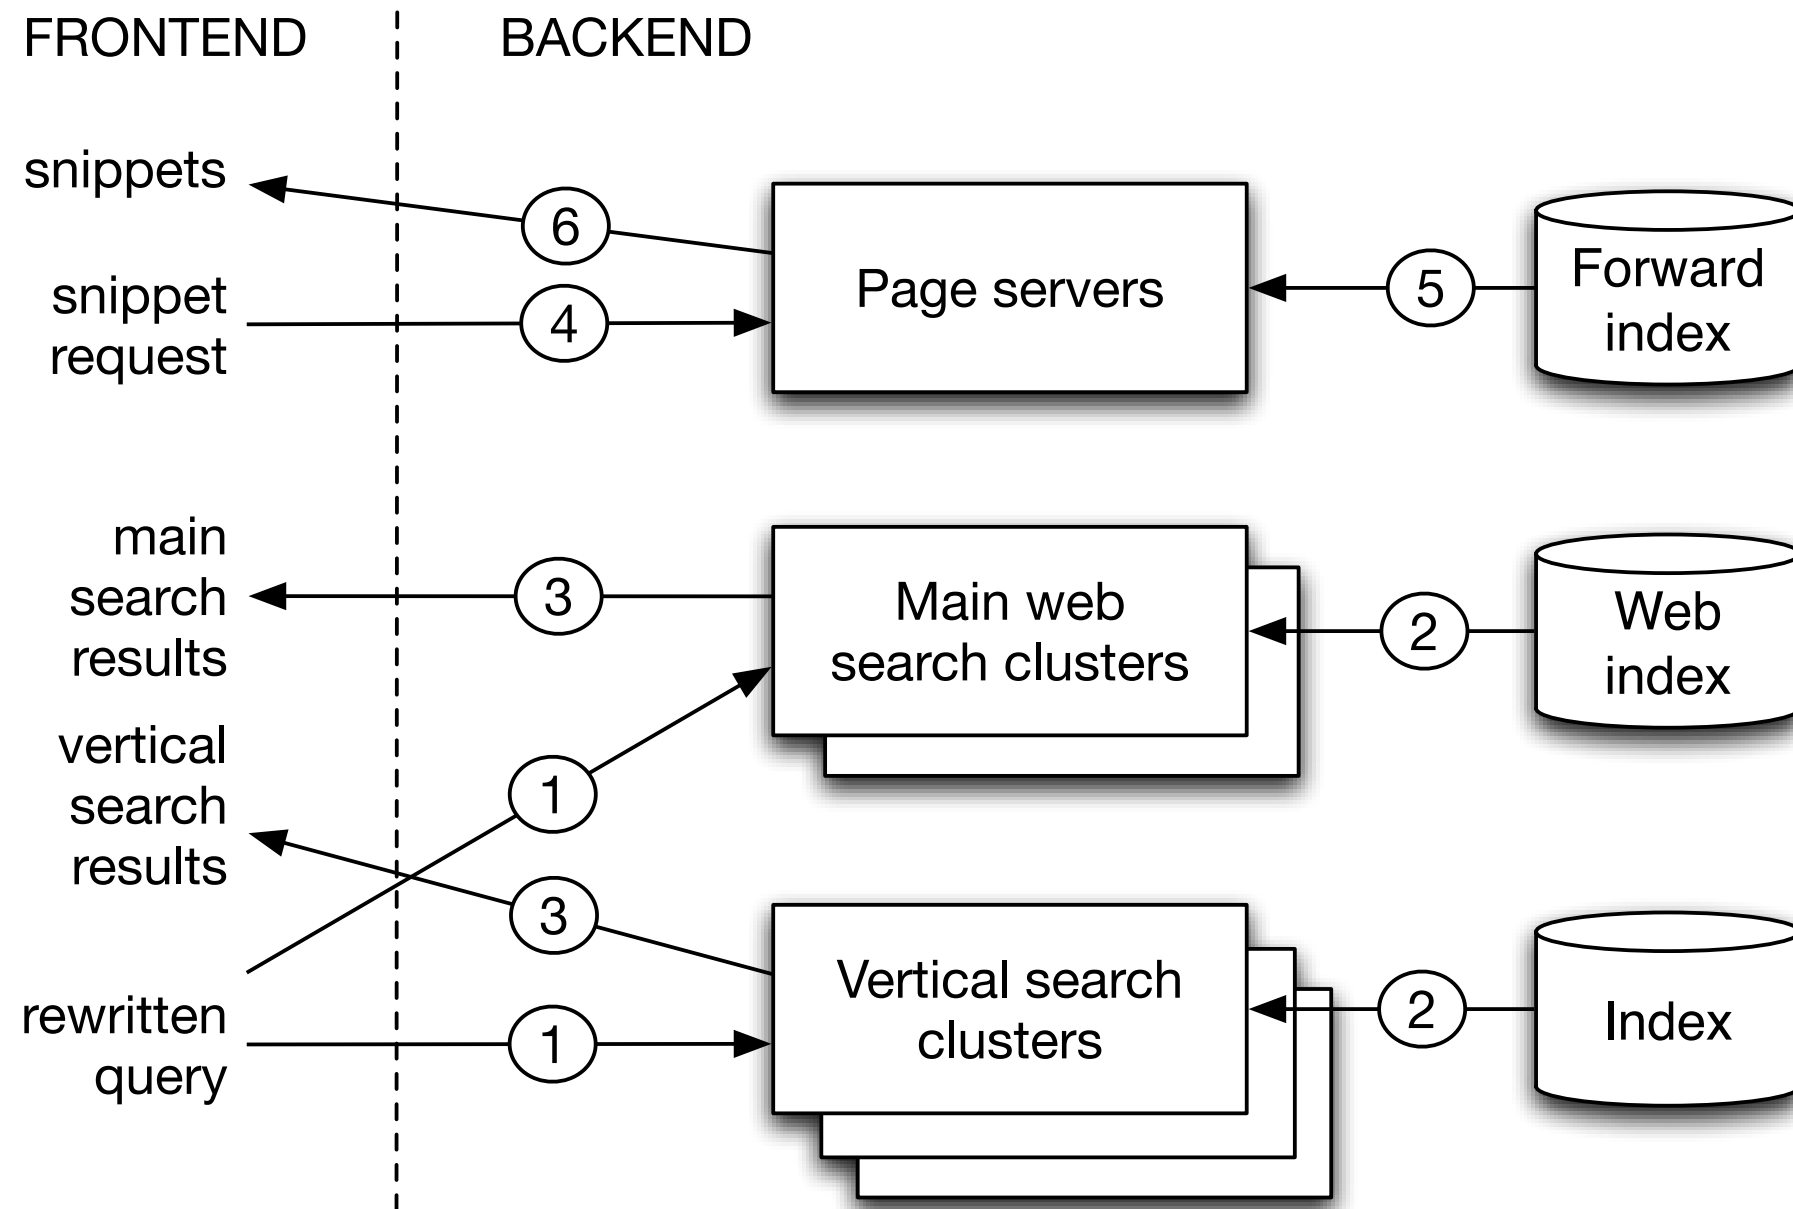
Query Interpretation System



Query Rewriting

- The original user query: “*amusement arcades in New York*”.
- Internal system query: AND(OR(PHRASE(*amusement arcade*),
PHRASE(*video arcade*)
)
LOCATION(*new york*)
)
- Applied modifications
 1. stop word “*in*” is removed
 2. term “*arcades*” is converted into its singular form “*arcade*”
 3. “*amusement arcade*” is detected as a phrase and expanded to “*video arcade*”
 4. “*New York*” is detected as a location and converted to lower case

Results Generation




Success Measures

- Quality measures
 - result quality: degree to which returned answers meet user's information need.
- Performance measures
 - latency: response time delay experienced by the user
 - peak throughput: number of queries that can be processed per unit of time without any degradation in other metrics

Measuring Results Quality

- Sources of feedback
 - editorial
 - user clicks



- Common measures
 - recall
 - precision
 - DCG
 - NDCG

	Ranking 1	Ranking 2	Optimal
	1. R	1. 	1. R
	2. 	2. R	2. R
	3. 	3. 	3. R
	4. 	4. 	4.
Recall:	1/3	1/3	1
Precision:	1/4	1/4	3/4
DCG:	1	0.63	$1+0.63+0.5=2.13$
NDCG:	$1/2.13$	$0.63/2.13$	

Query Evaluation

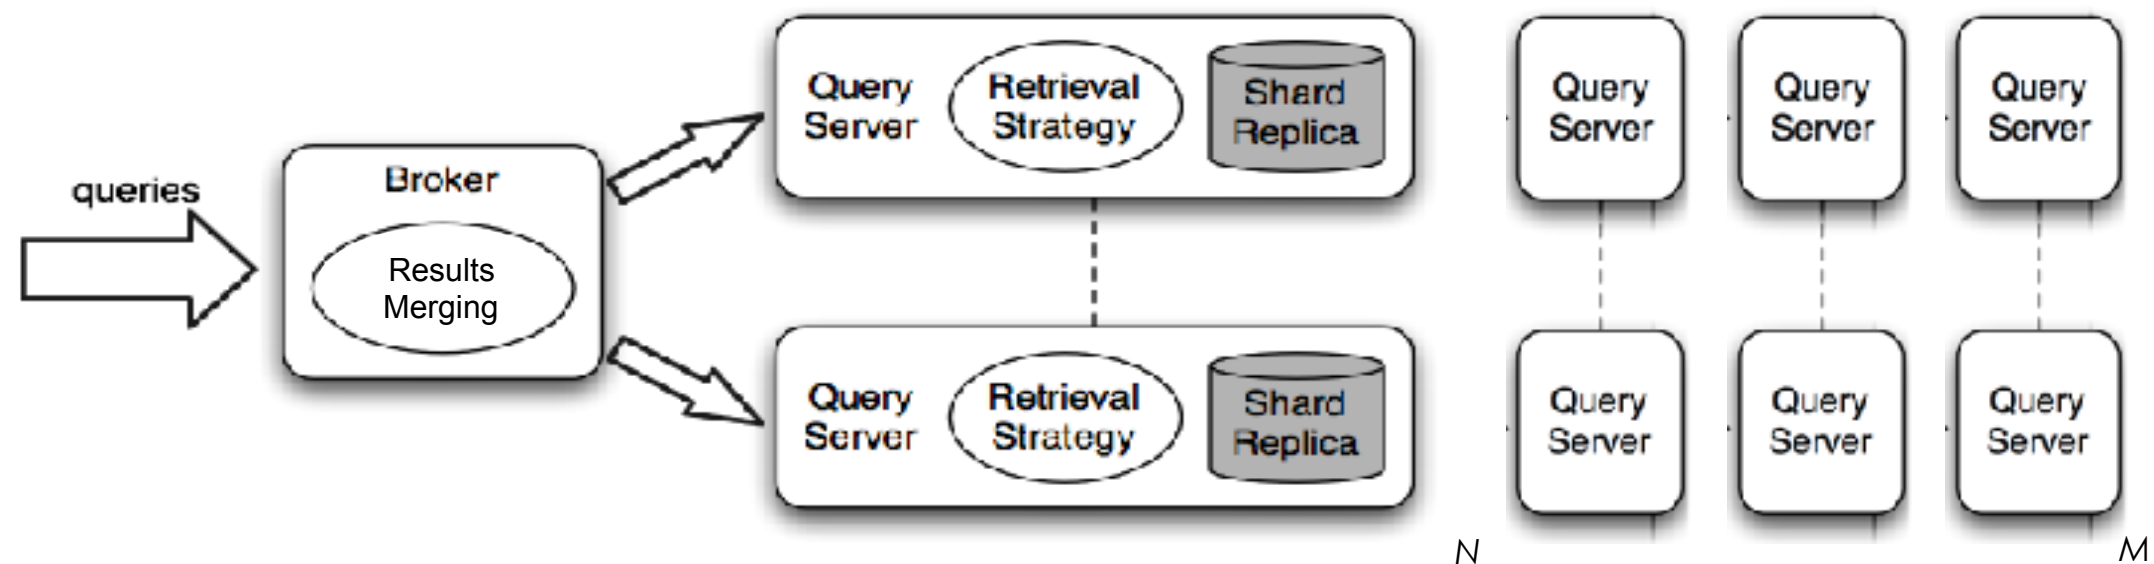
- Conjunctive (AND) vs. Disjunctive (OR)
- Boolean retrieval
 - Suitable for small/medium collections
 - Returns all documents matching the query
- Ranked retrieval
 - Suitable for Web-scale collections
 - Requires a similarity function between queries and documents

$$\text{SCORE}_q(d) = \sum_{t \in q} s_t(q, d)$$

- Returns only the top k document

Query Processing Infrastructures

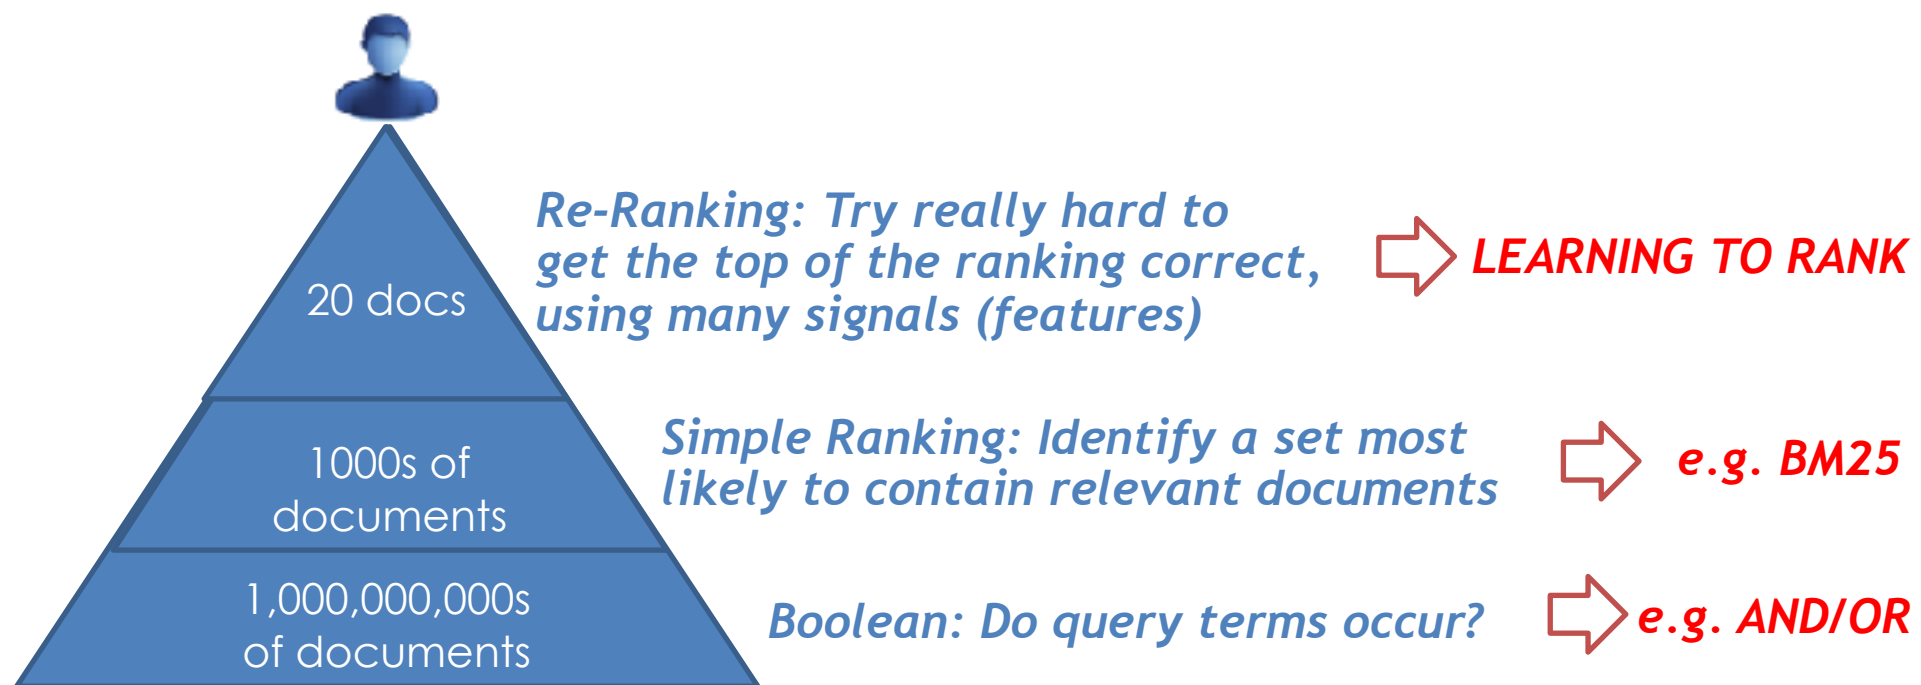
- To achieve efficiency at Big Data scale, search engines use many servers:



- N & M can be very big:
 - Microsoft's Bing search engine has "*hundreds of thousands of query servers*"

Ranking Cascades


- Typically, in web-scale search, the ranking process can be conceptually seen as a series of cascades [1]
 - Rank some documents
 - Pass top-ranked onto next cascade for refined re-ranking



Query Evaluation

- Normal strategies make a pass on the postings lists for each query term
 - This can be done **Term-at-a-Time (TAAT)** – one query term at a time
 - Or **Document-at-a-time (DAAT)** – all query terms in parallel
- We will explain these, before showing how we can improve them

Term-at-a-Time (TAAT)




Advantages:

- Simple

Disadvantages:

- Requires lots of memory to contain partial scores for all documents
- Difficult to do Boolean or phrase queries, as we don't have a document's postings for all query terms at once

Document-at-a-Time (DAAT)



Advantages:

- Reduced memory compared to TAAT (and hence faster)
- Supports Boolean query operators, phrases, etc.

Disadvantages:

- Slightly more complex to implement

Most commercial search engines are reported to use DAAT

TAAT vs DAAT

- TAAT and DAAT have been the cornerstones of query evaluation in IR systems since 1970s.
- The plain implementations of these two strategies are seldom used anymore, since many optimizations have been proposed during the years
- Several known systems in production today, from large scale search engines as Google and Yahoo!, to open source text indexing packages as Lucene and Lemur, use some optimized variation of these strategies
- [Turtle and Flood, 1995] were the first to argue that DAAT could beat TAAT in practical environments.
 - For large corpora, DAAT implementations require a smaller run-time memory footprint
- [Fontoura et al., 2011] report experiments on small (200k docs) and large indexes (3M docs), with short (4.26 terms) and long (57.76 terms) queries (times are in microseconds):

Small index		
	Short queries	Long queries
TAAT	141.0	1,694.6
DAAT	193.0	4,554.6
Large index		
	Short queries	Long queries
TAAT	3,777.6	18,913.0
DAAT	3,581.3	26,778.3

Dynamic Pruning

- Dynamic pruning strategies aim to make scoring faster by only **scoring a subset** of the documents
 - The core assumption of these approaches is that the user is only interested in the top K results, say $K=20$
 - During query scoring, it is possible to determine if a document cannot make the top K ranked results
 - Hence, the scoring of such documents can be terminated early, or skipped entirely, without damaging retrieval effectiveness to rank K
- We call this “*safe-to-rank K*”
- Dynamic pruning is based upon
 - *Early termination*
 - *Comparing upper bounds on retrieval scores with thresholds*

Techniques

- MaxScore
 - **Early termination**: does not compute scores for documents that won't be retrieved by comparing **upper bounds** with a score **threshold**
- WAND
 - **Approximate evaluation**: does not consider documents with approximate scores (sum of **upper bounds**) lower than **threshold**
 - Therefore, it focuses on the combinations of terms needed (**wAND**)
- BlockMaxWand
 - SOTA variant of WAND that uses benefits from the block-layout of posting lists

MaxScore Example

Sorted in increasing order of
term upper bound

term
upper
bounds



partial
score



posting list
iterator




current
docid

non-essential lists	1.3	ub[0]	1.3	p[0]	11	13	23
	2.1	ub[1]	3.4	p[1]	11	22	26
essential lists	3.7	ub[2]	8.1	p[2]	11	24	27
	4.0	ub[3]	11.1	p[3]	11	23	28
threshold		6.0	0.0				

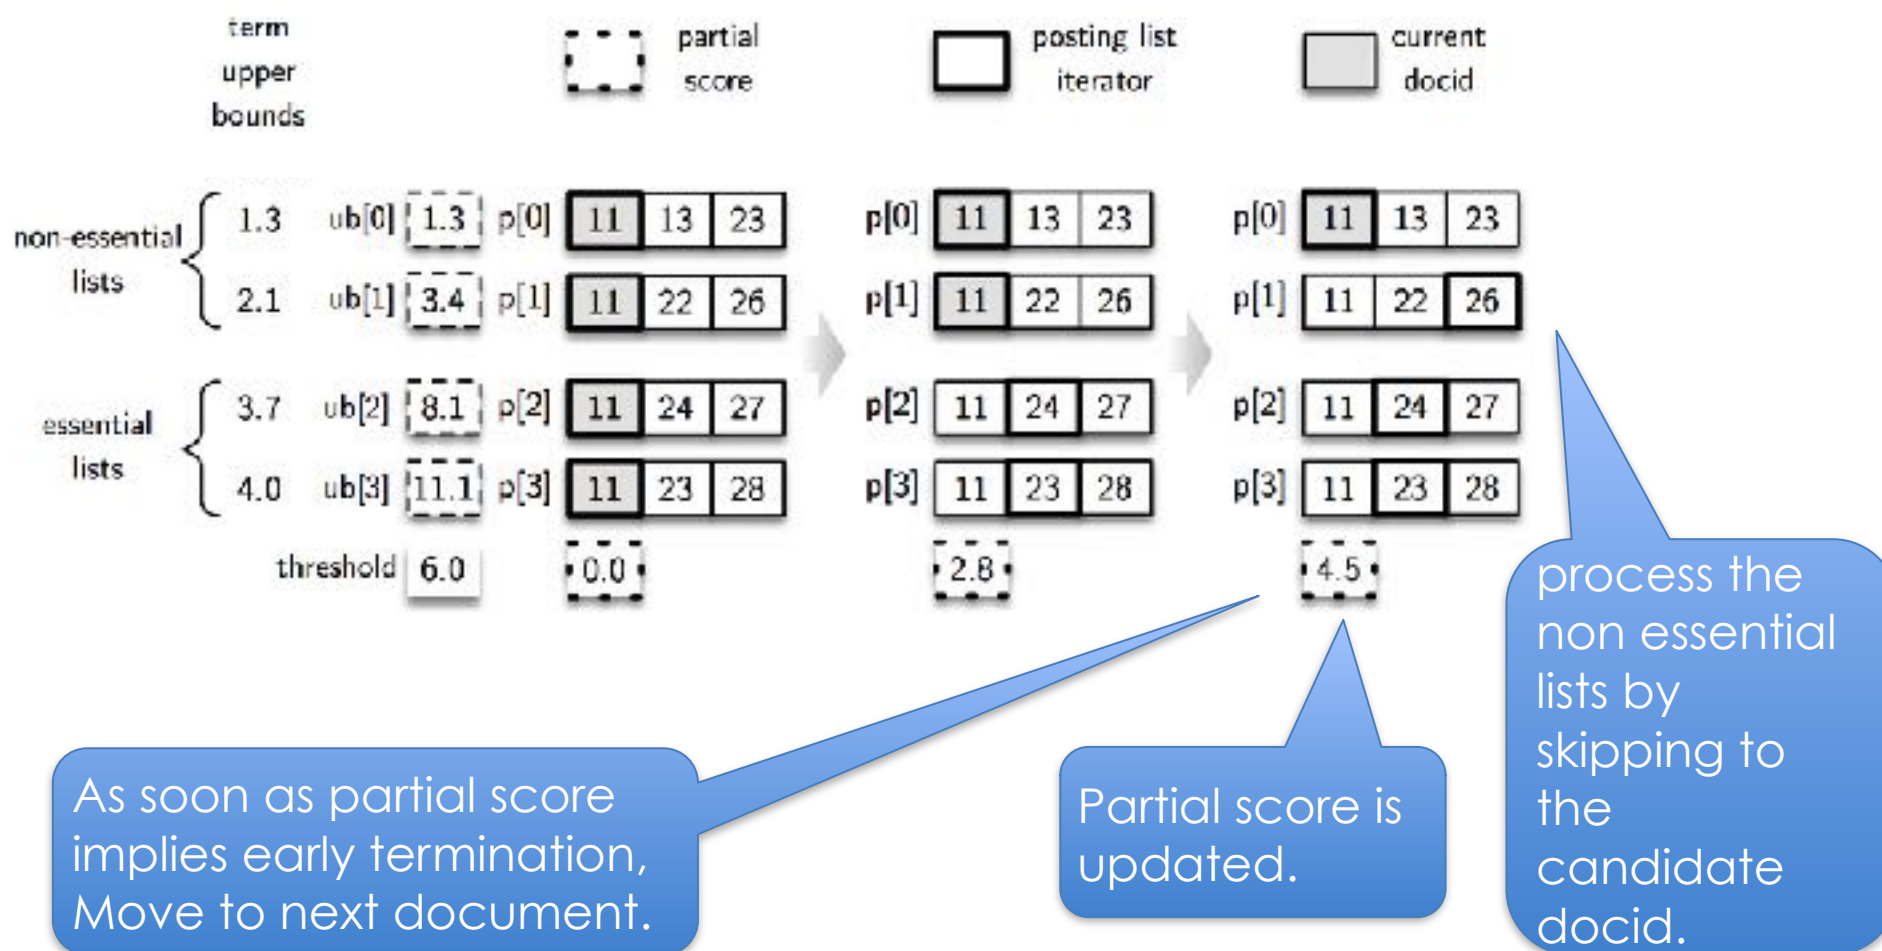
no document can be returned as a
top k results if it appears in the non-
essential lists only, i.e., $1.3 + 3.4 < 6.0$

MaxScore Example



Essential lists are processed in DAAT, and their score contributions computed

MaxScore Example



Learning

- How to choose term weighting models?
 - Different term weighting models have different **assumptions** about how relevant documents should be retrieved
- Also:
 - **Field-based models**: term occurrences in different fields matter differently
 - **Proximity-models**: close co-occurrences matter more
 - **Priors**: documents with particular lengths or URL/inlink values matter more
 - **Query Features**: Long queries, difficult queries, query type

How to combine all these easily and appropriately?

Learning to Rank

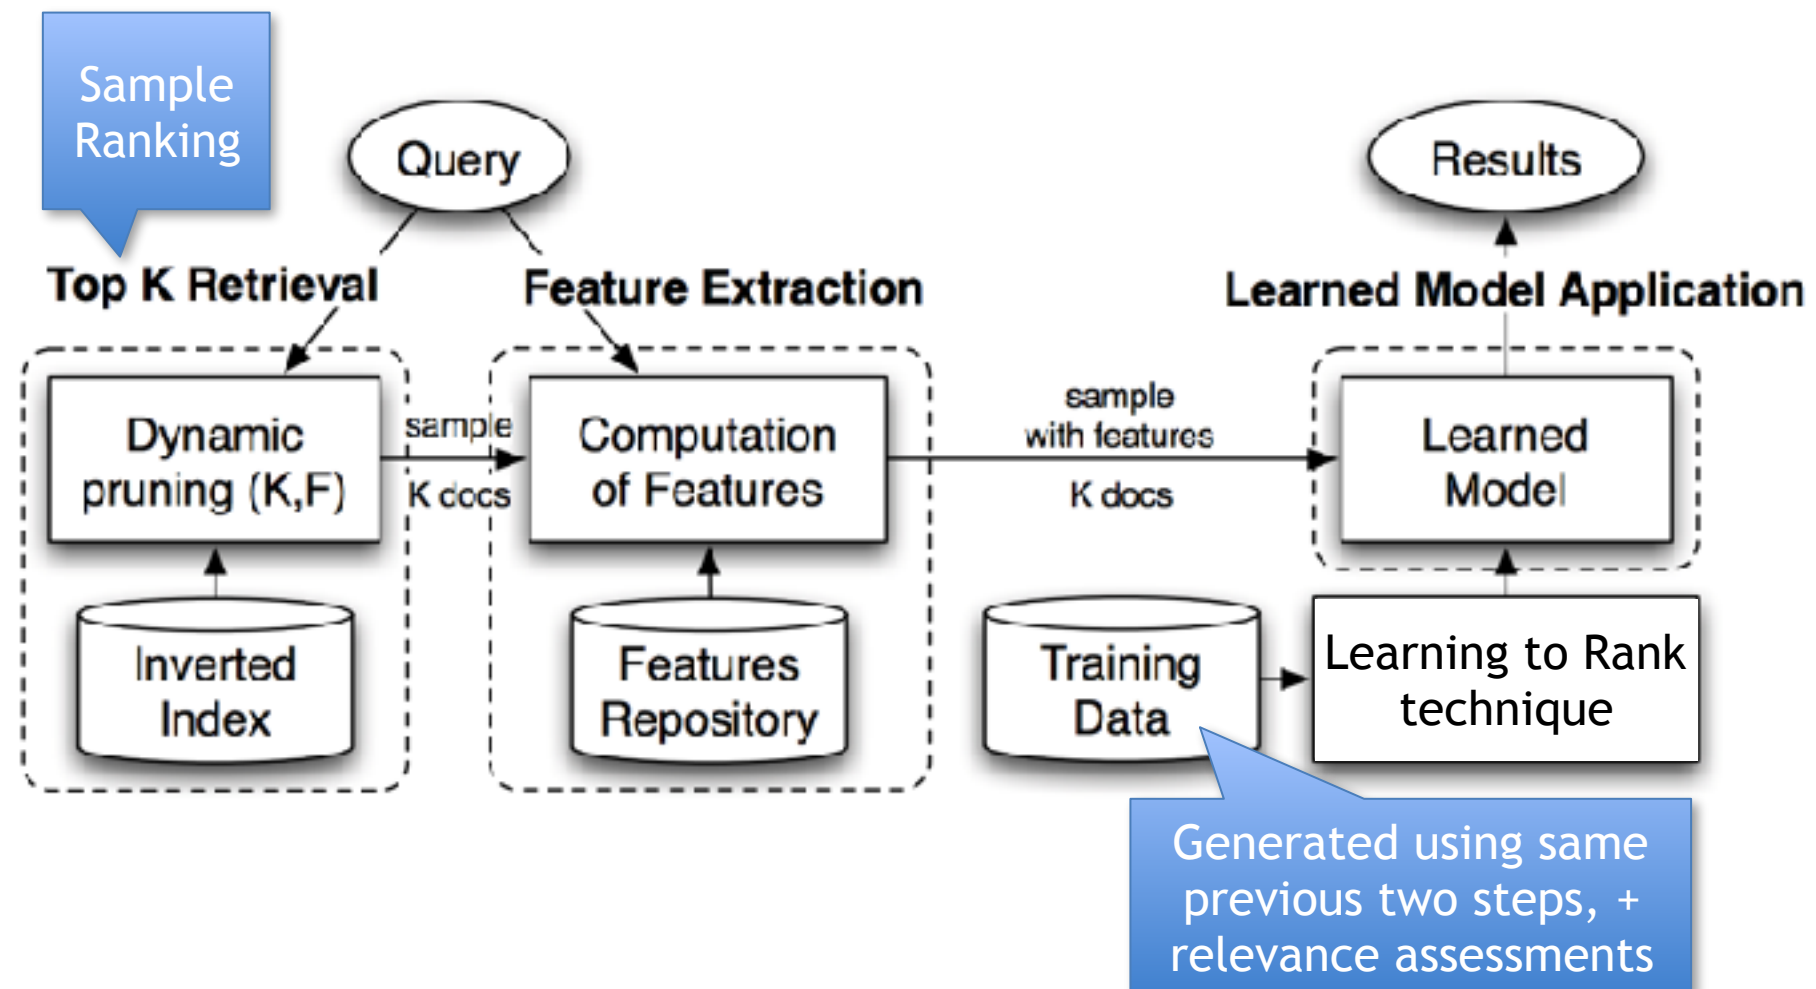
- Application of tailored machine learning techniques to automatically (select and) weight retrieval *features*
 - Based on **training data** with relevance assessments
- Recently, learning to rank has been popularised by commercial search engines (e.g. Bing, Baidu, Yandex)
 - They require large **training datasets**, possibly instantiated from **click-through data**
 - **Click-through data** has facilitated the deployment of learning approaches

Features

Typically, commercial search engines use hundreds of features for ranking documents, usually categorised as follows:

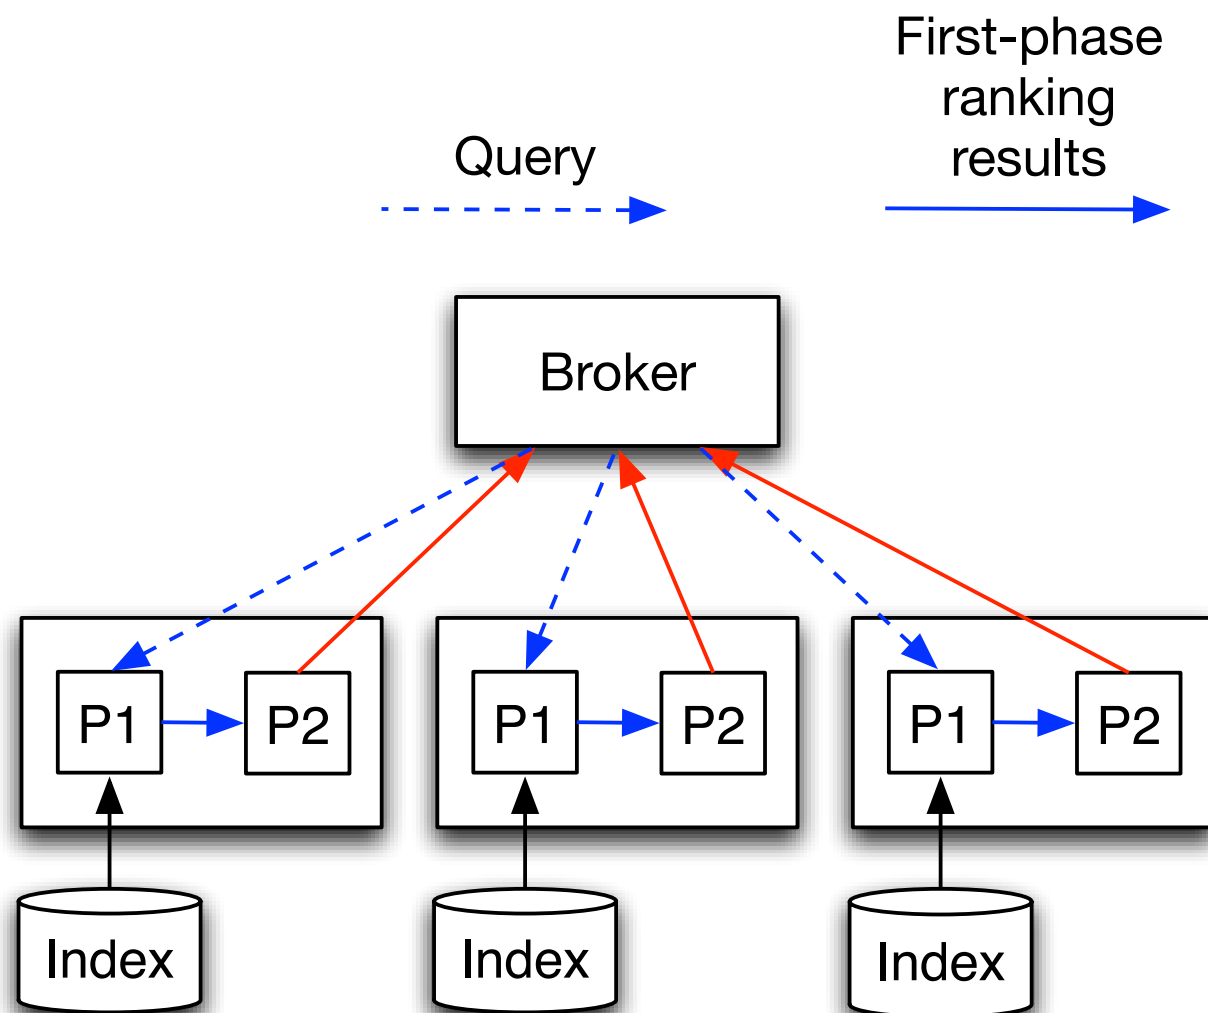
Name	Varies depending on...		Examples
	Query	Document	
Query Dependent Features	✓	✓	Weighting models, e.g. BM25, PL2 Proximity models, e.g. Markov Random Fields Field-based weighting models, e.g. PL2F
Query Independent Features	x	✓	PageRank, number of inlinks Spamminess
Query Features	✓	x	Query length Presence of entities

Schematically

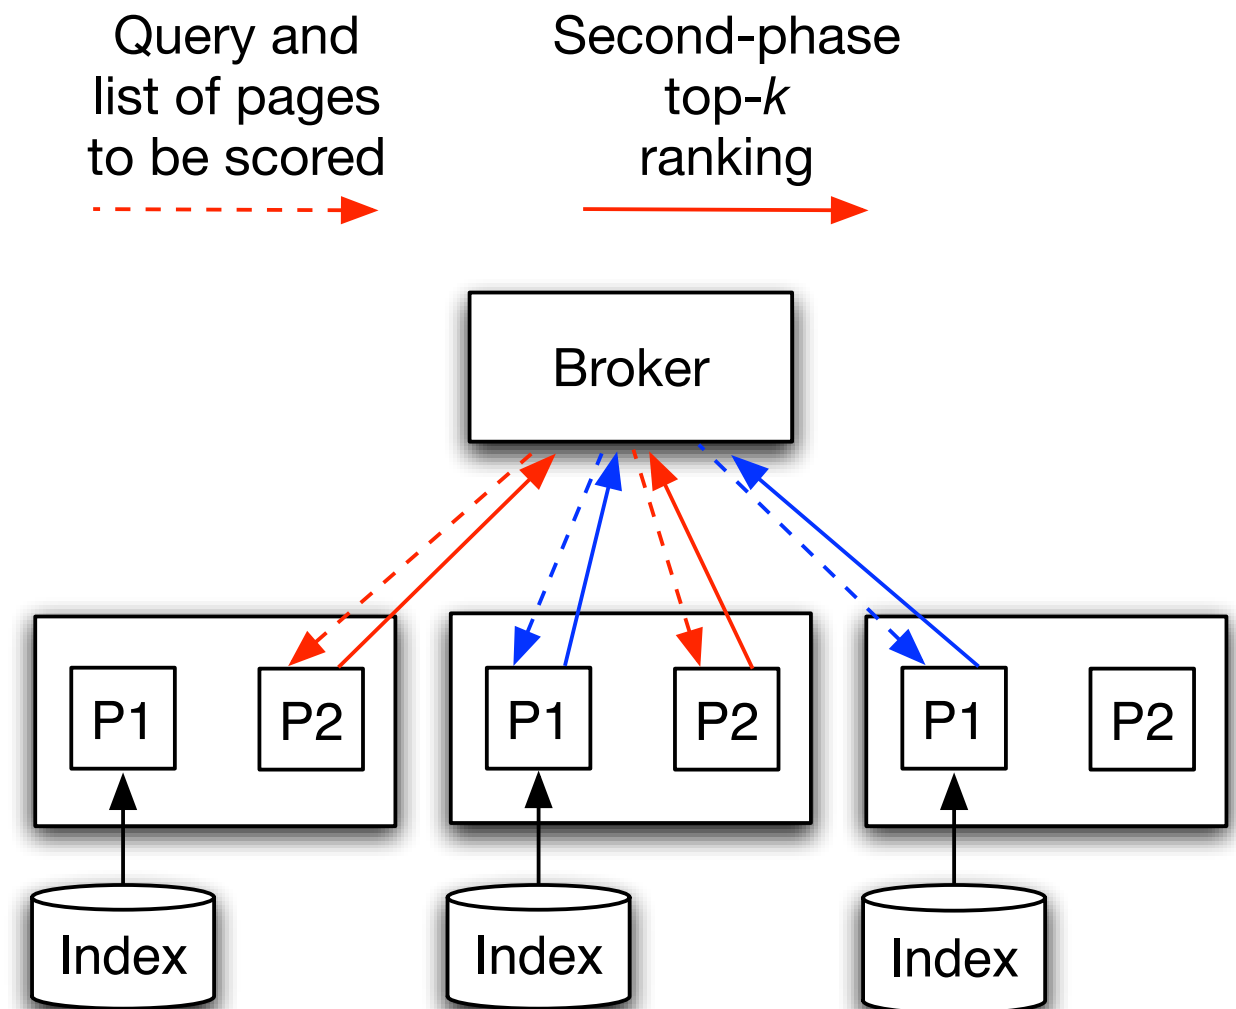


Parallel Query Processing on a search cluster

Document-based partitioning



Term-based partitioning



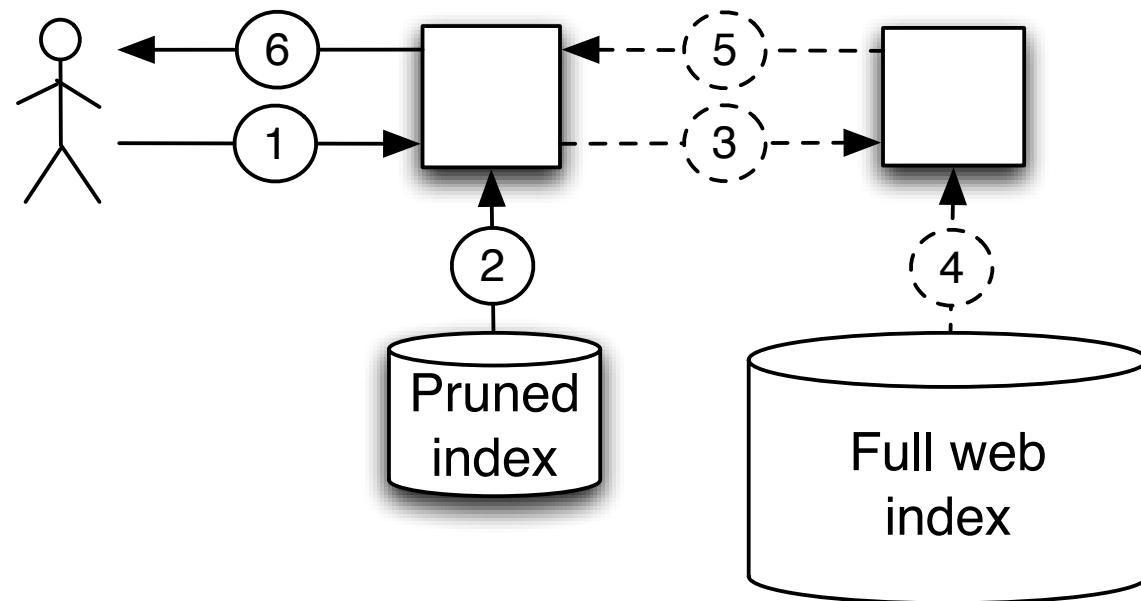
Architectural Optimizations

- Static index pruning
- Index tiering
- Selective Search
- Caching

Static Index Pruning

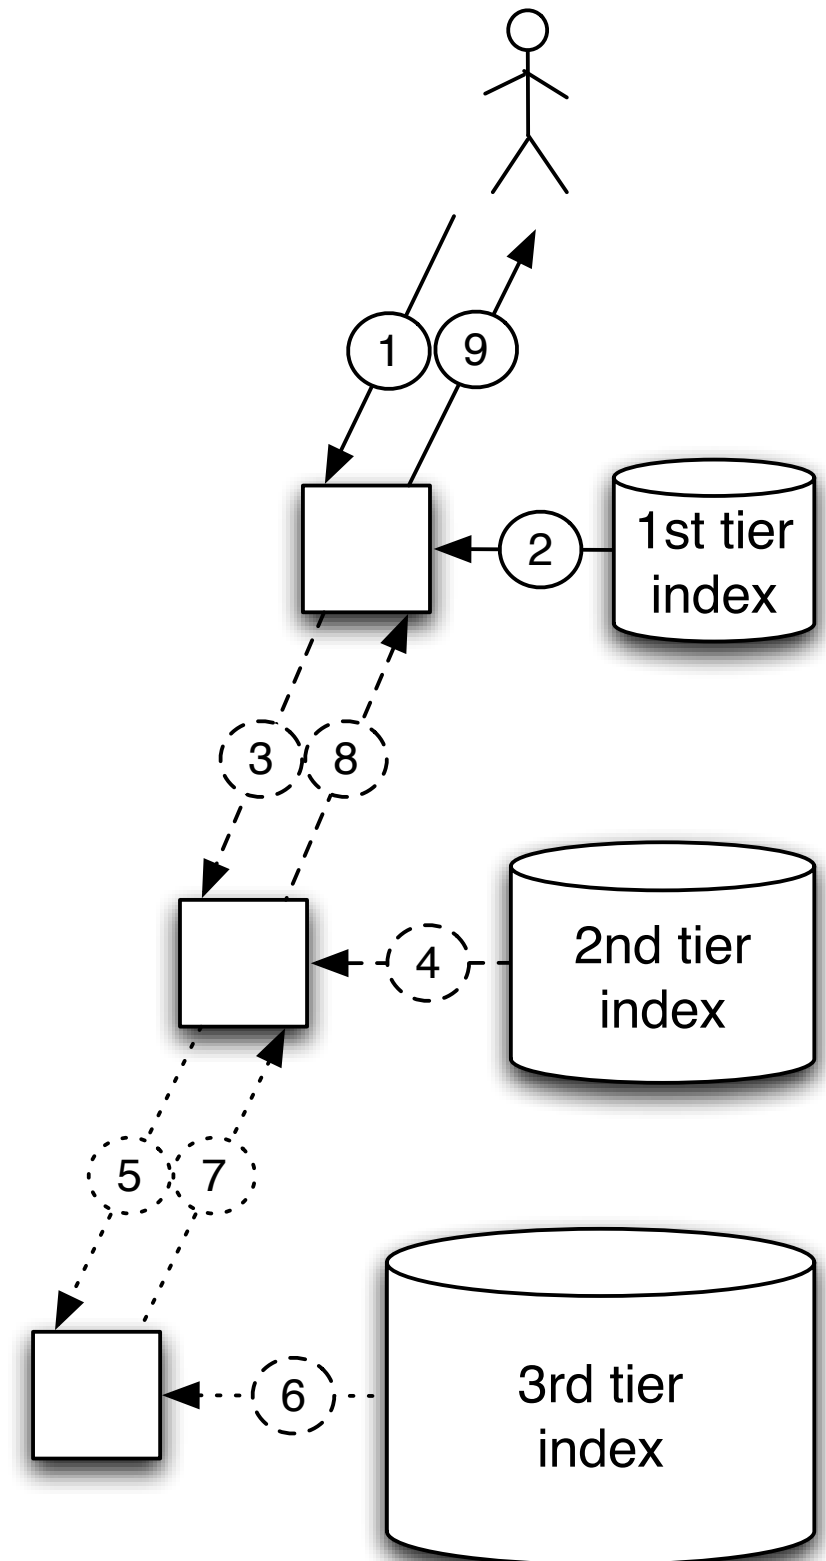
- Idea: to create a small version of the search index that can accurately answer most search queries

- Techniques
 - term-based pruning
 - doc-based pruning
- Result quality
 - guaranteed
 - not guaranteed




Tiering

- A sequence of sub-indexes
 - former sub-indexes are small and keep more important documents
 - later sub-indexes are larger and keep less important documents
 - a query is processed selectively only on the first n tiers
- Two decisions need to be made
 - tiering (offline): how to place documents in different tiers
 - fall-through (online): at which tier to stop processing the query



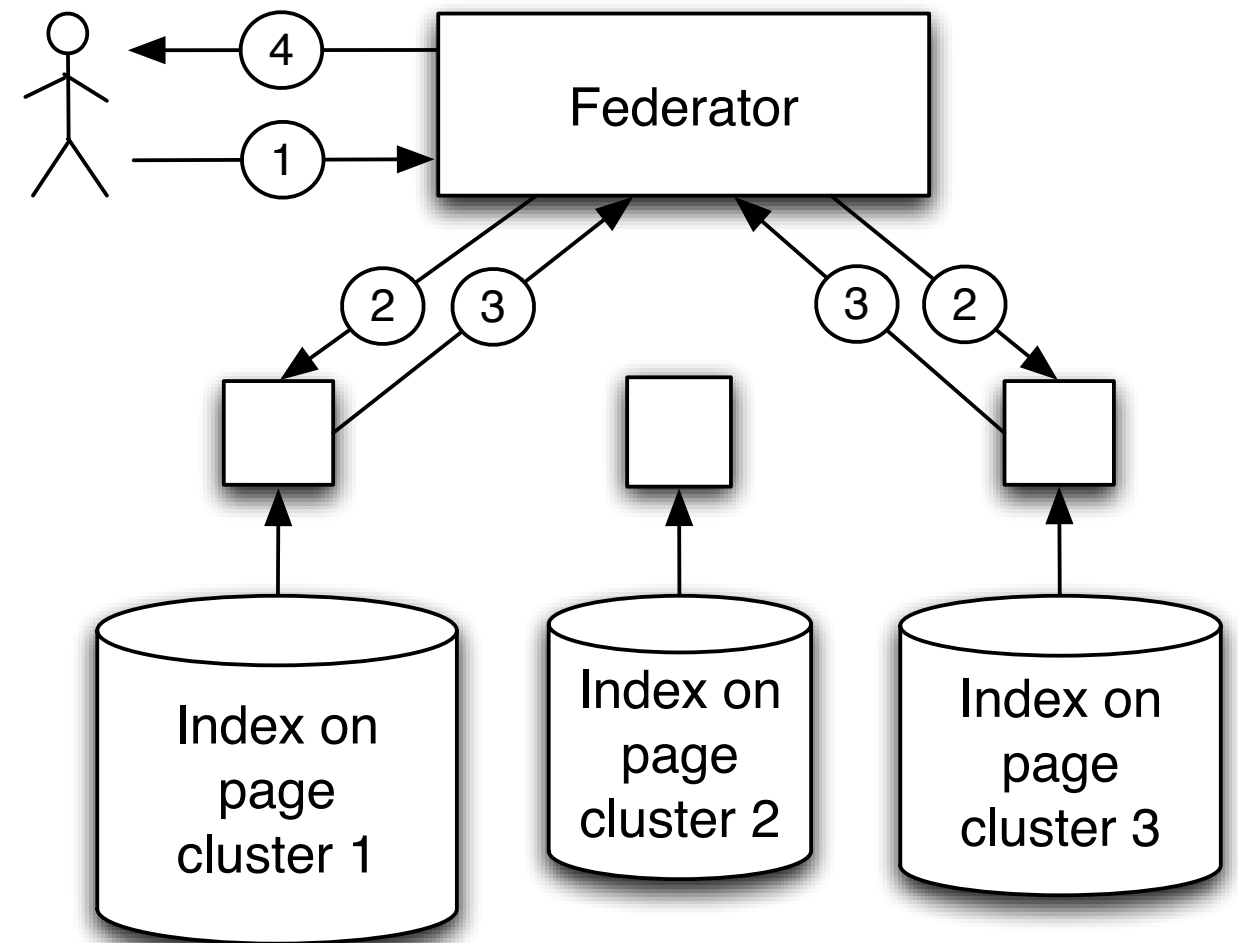
Tiering

- Tiering strategy is based on some document importance metric
 - PageRank
 - click count
 - spam score
- Fall-through strategy
 - query the next index until there are enough results
 - query the next index until search result quality is good
 - predict the next tier's result quality by machine learning



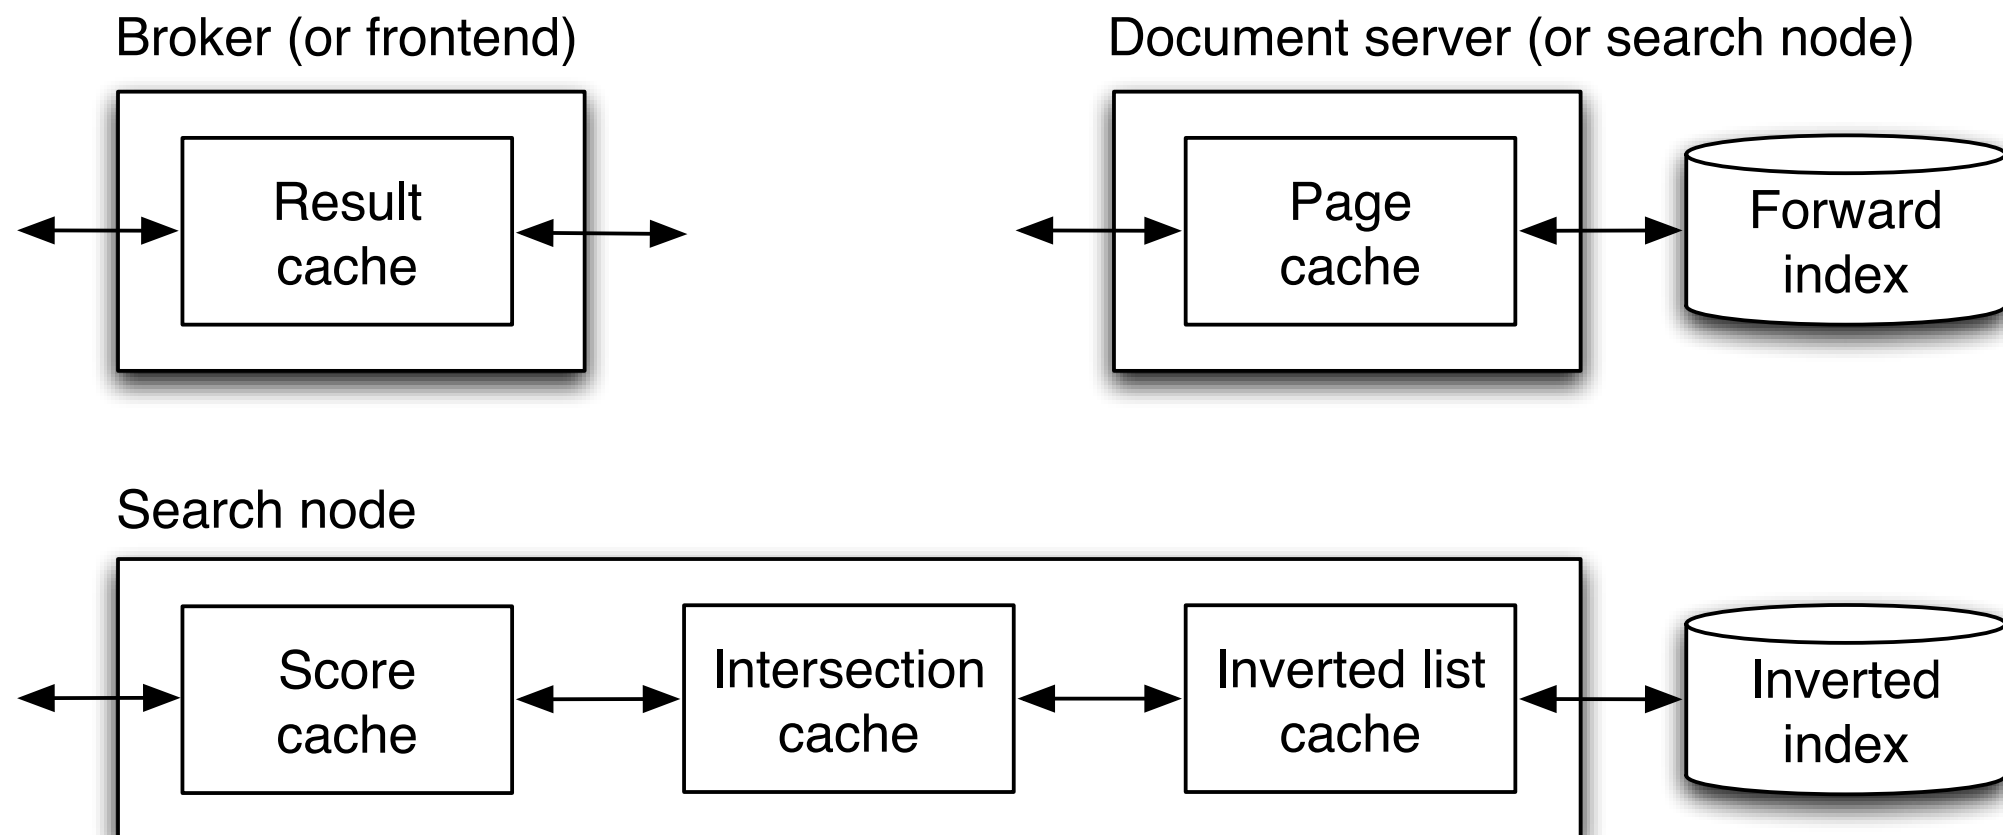
Selective Search

- Documents are clustered and a separate index is built
 - similarity between documents
 - co-click likelihood
- A query is processed on the indexes associated with the most similar n clusters
- Reduces the workload, but suffers from the load imbalance problem
 - query topic distribution may be skewed
 - certain indexes have to be queried much more often




Caching

Main caches in search engines: result cache, score cache, intersection cache, inverted list cache, page cache



Partitioned Search Architectures

- Key points
 - multiple, regional data centers (sites)
 - user-to-center assignment
 - partitioned web index
 - partial document replication
- Enables
 - local web crawling
 - query processing with selective forwarding



Open Source Search Engines

- ATIRE: Search engine (BSD license)
- DataparkSearch: Website search engine (GNU GPL)
- Elasticsearch: Search server based on Lucene (Apache license 2.0)
- Galago: Search engine toolkit (BSD license)
- Indri: Search engine (BSD-inspired license)
- **Lucene**: Search engine library (Apache license 2.0)
- **MG4J**: Distributed search engine (GNU LGPL)
- mnoGoSearch: Website search engine (GNU GPL)
- **Solr**: Distributed search engine based on Lucene (Apache license 2.0)
- Seeks: User-centric search engine (Affero GPLv3)
- Sphinx: Search engine (GNU GPLv2)
- **Terrier**: Search engine (Mozilla public license)
- Wumpus: Desktop search engine (GNU GPL)
- Zettair: Search engine (BSD license)

Open Source Search Engines

- ATIRE: Search engine (BSD license)
- DataparkSearch: Website search engine (GNU GPL)
- Elasticsearch: Search server based on Lucene (Apache license 2.0)
- Galago: Search engine toolkit (BSD license)
- Indri: Search engine (BSD-inspired license)
- Lucene: Search engine library (Apache license 2.0)
- MG4J: Distributed search engine (GNU LGPL)
- mnoGoSearch: Website search engine (GNU GPL)
- Solr: Distributed search engine based on Lucene (Apache license 2.0)
- Seeks: User-centric search engine (Affero GPLv3)
- Sphinx: Search engine (GNU GPLv2)
- Terrier: Search engine (Mozilla public license)
- Wumpus: Desktop search engine (GNU GPL)
- Zettair: Search engine (BSD license)